

Neural Branch Predictors

Quinn Pham

You are at the LRT station chatting with a fellow student. She asks, what is the big deal with **neural branch prediction**? Why did Quinn spend a lecture talking about it? Your train is coming, you can only say a few sentences. What would you tell her?

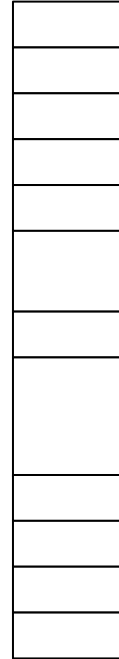
Two-level branch predictors

Two-level branch predictors

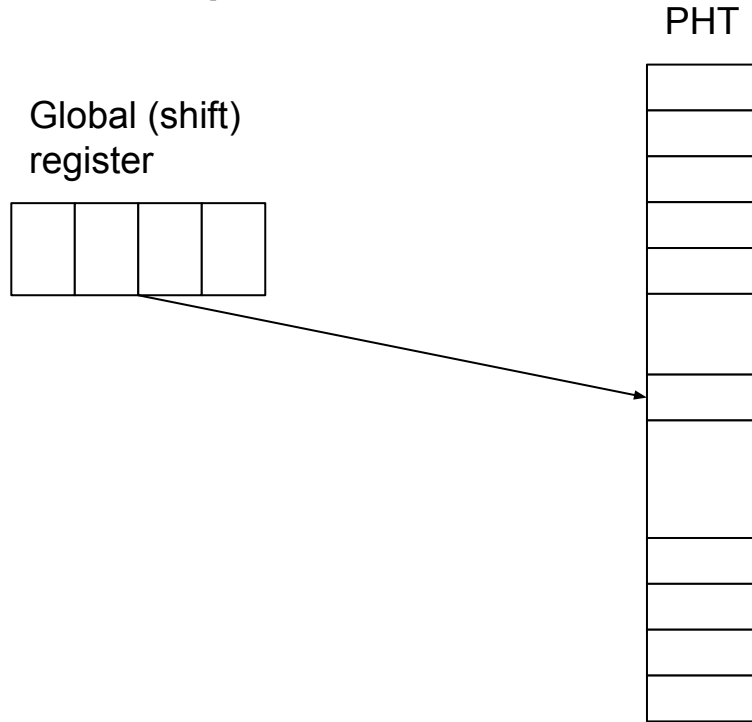
Global (shift)
register



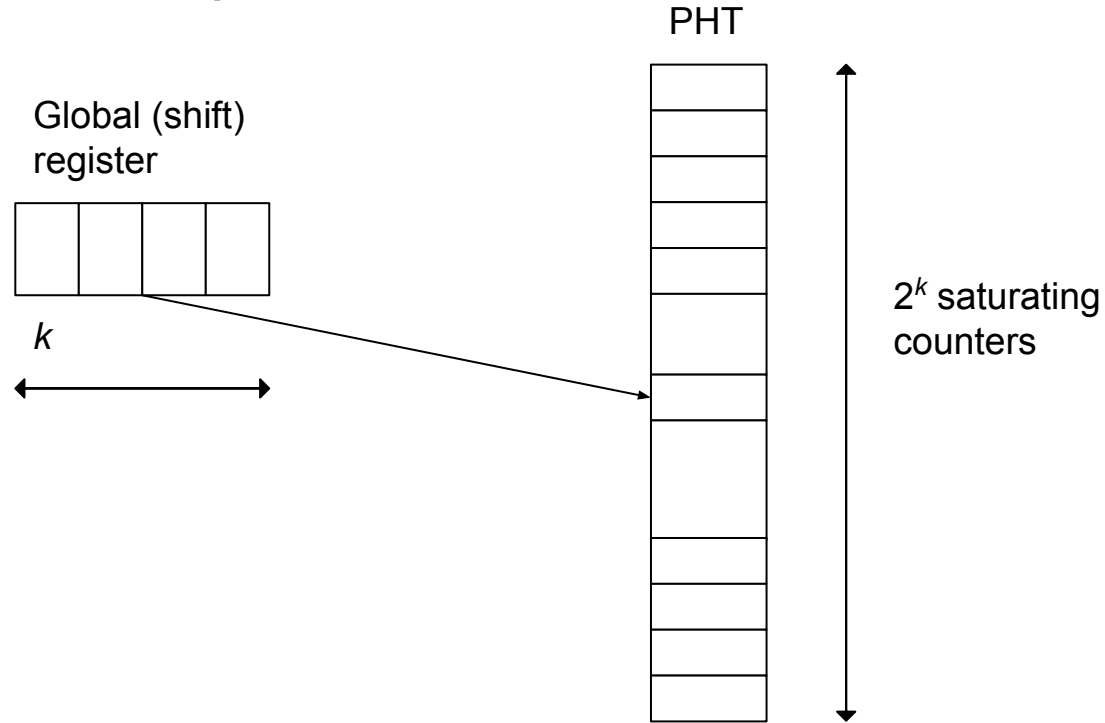
PHT



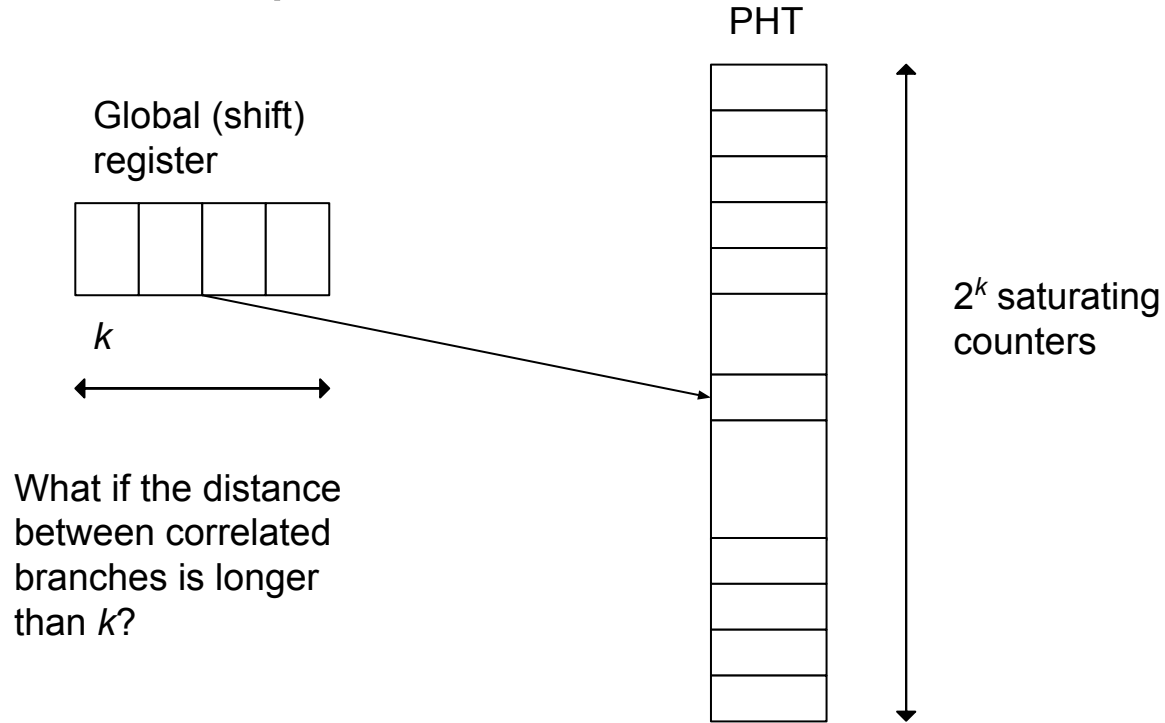
Two-level branch predictors



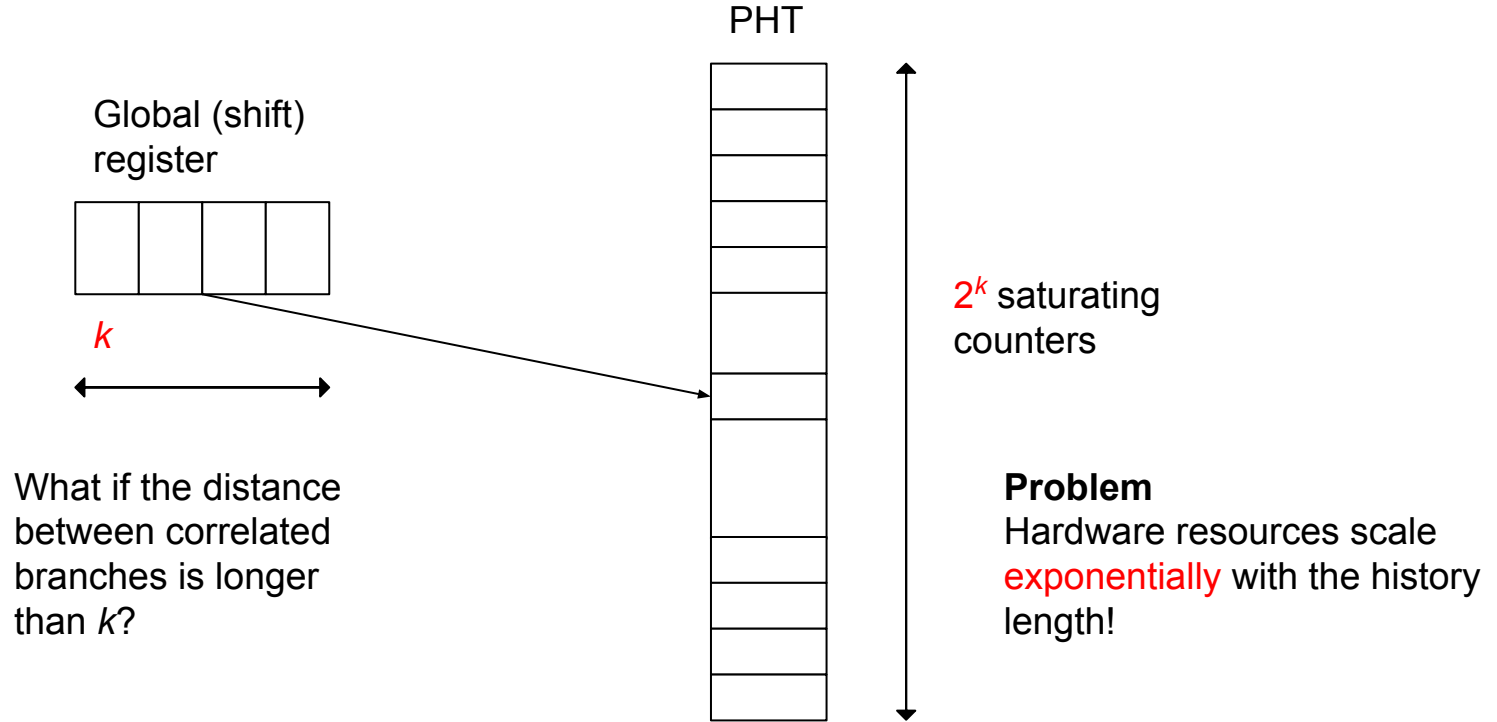
Two-level branch predictors



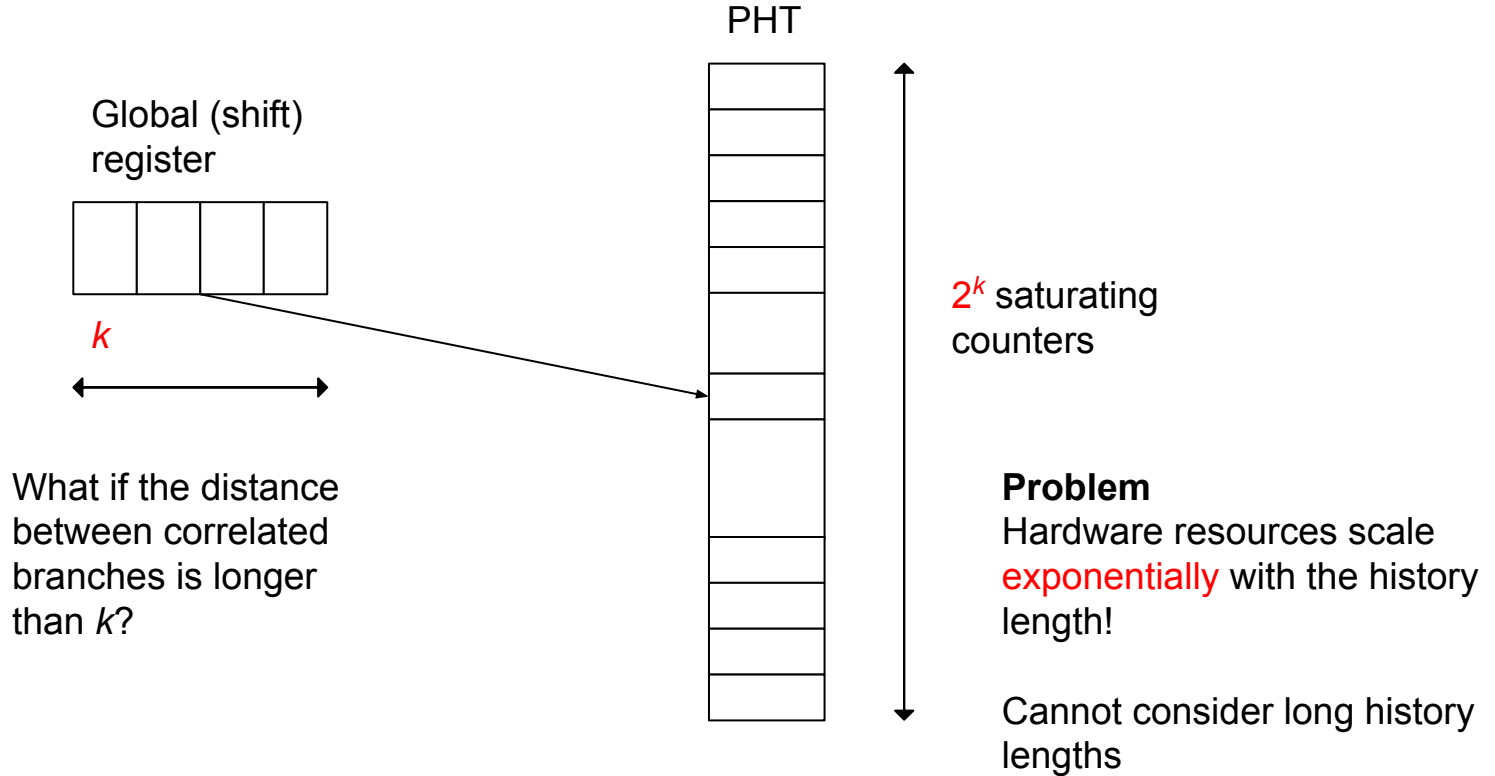
Two-level branch predictors



Two-level branch predictors



Two-level branch predictors



Dynamic Branch Prediction with Perceptrons

Daniel A. Jiménez Calvin Lin

Department of Computer Sciences

The University of Texas at Austin

Austin, TX 78712

`{ djimenez, lin}@cs.utexas.edu`



Dynamic Branch Prediction with Perceptrons

Daniel A. Jiménez Calvin Lin

Department of Computer Sciences

The University of Texas at Austin

Austin, TX 78712

`{ djimenez, lin}@cs.utexas.edu`



Replace the pattern history table (PHT) in the two-level predictor with a table of perceptrons

Dynamic Branch Prediction with Perceptrons

Daniel A. Jiménez Calvin Lin

Department of Computer Sciences

The University of Texas at Austin

Austin, TX 78712

`{ djimenez, lin}@cs.utexas.edu`



Replace the pattern history table (PHT) in the two-level predictor with a table of **perceptrons**

Perceptrons

Perceptrons

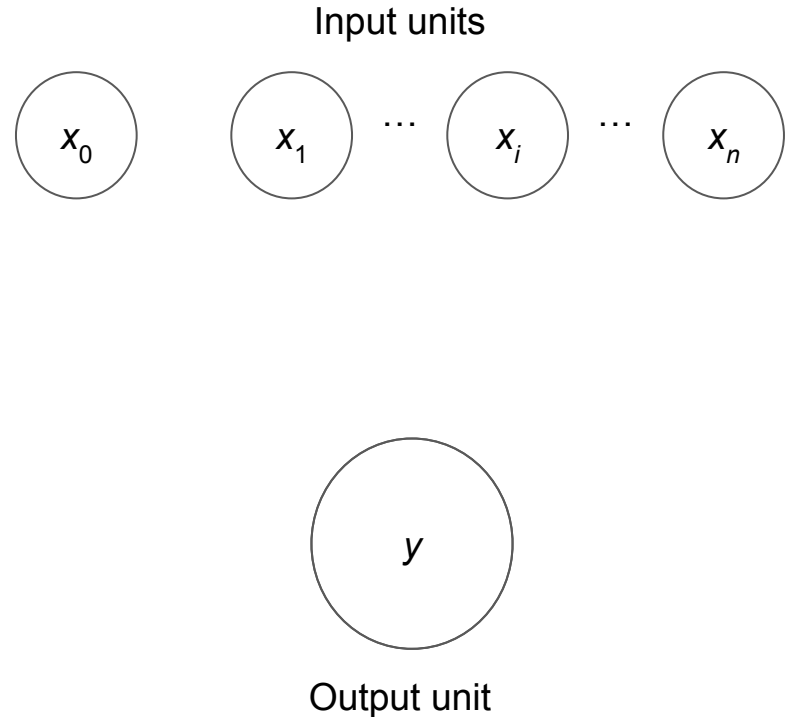
- Neural networks learn to compute a function using example inputs and outputs

Perceptrons

- Neural networks learn to compute a function using example inputs and outputs
- Single-layer neural network

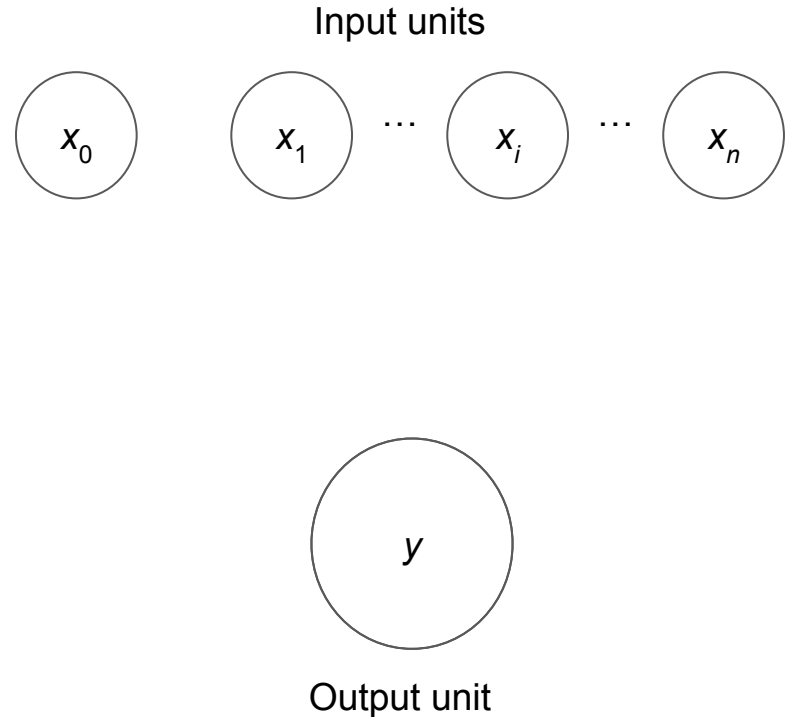
Perceptrons

- Neural networks learn to compute a function using example inputs and outputs
- Single-layer neural network



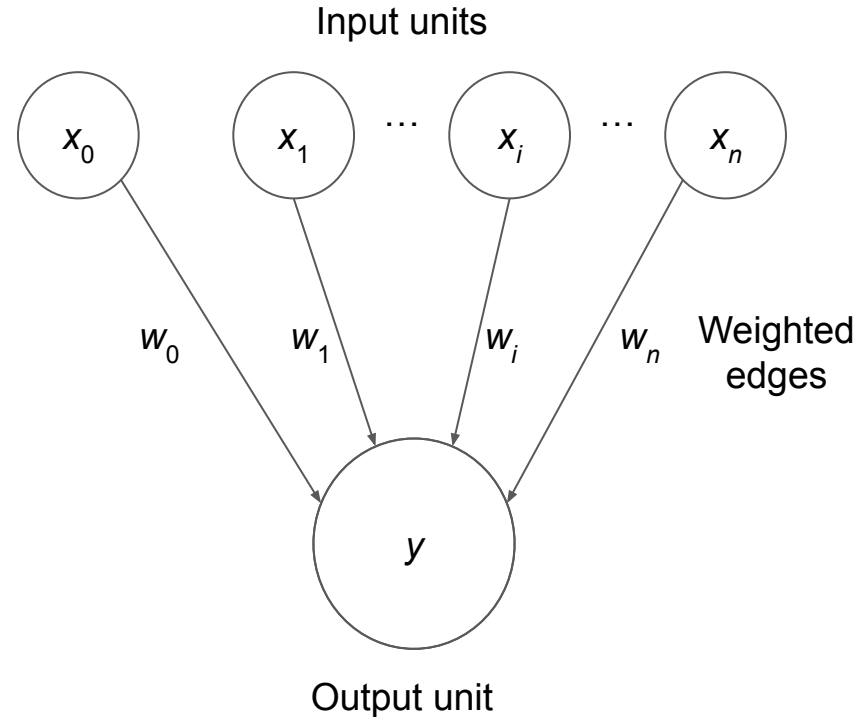
Perceptrons

- Neural networks learn to compute a function using example inputs and outputs
- Single-layer neural network
- Data is fed into input unit units and propagated through the network to the output unit



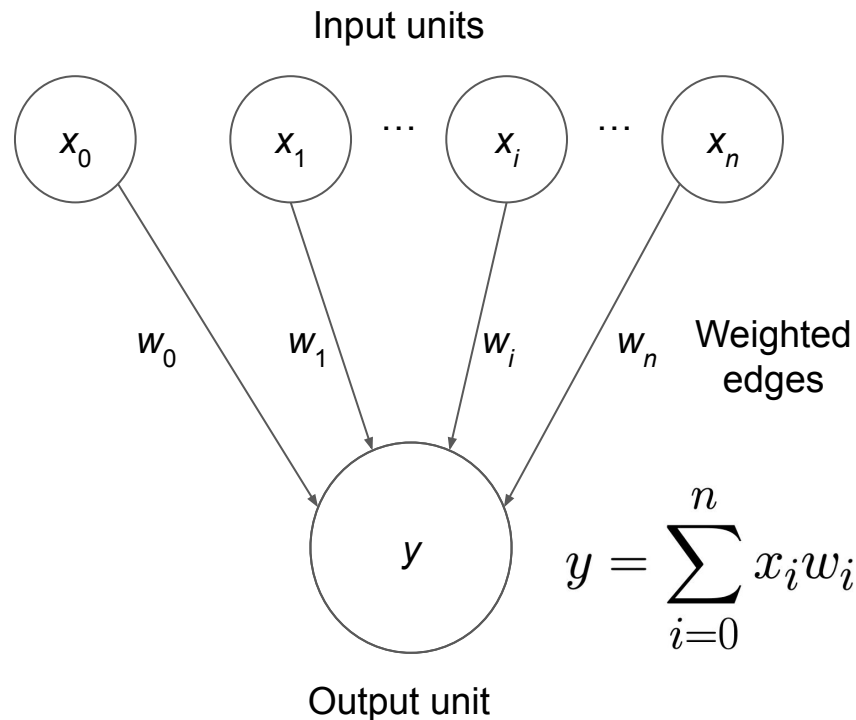
Perceptrons

- Neural networks learn to compute a function using example inputs and outputs
- Single-layer neural network
- Data is fed into input unit units and propagated through the network to the output unit



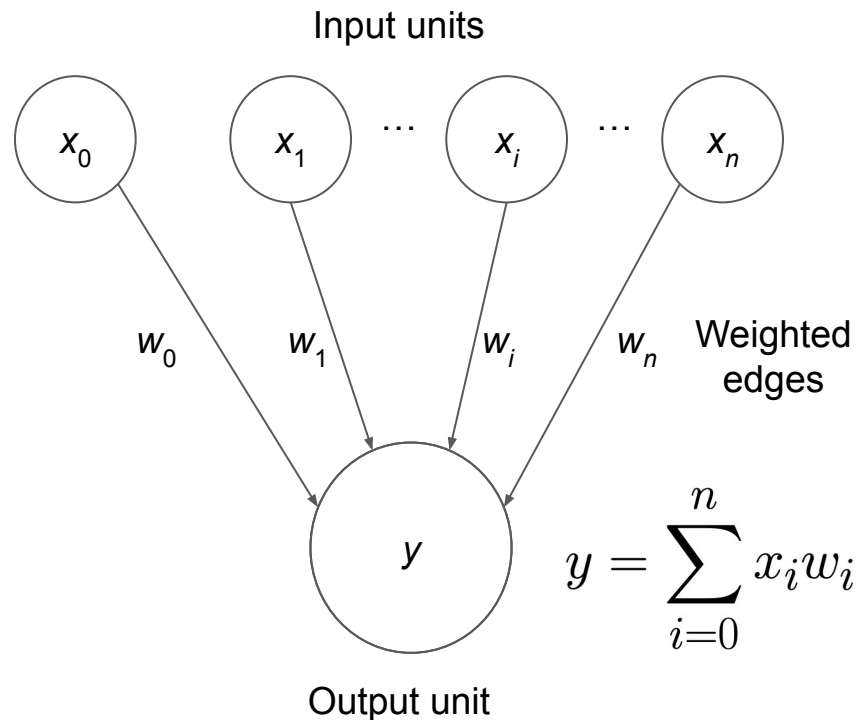
Perceptrons

- Neural networks learn to compute a function using example inputs and outputs
- Single-layer neural network
- Data is fed into input unit units and propagated through the network to the output unit

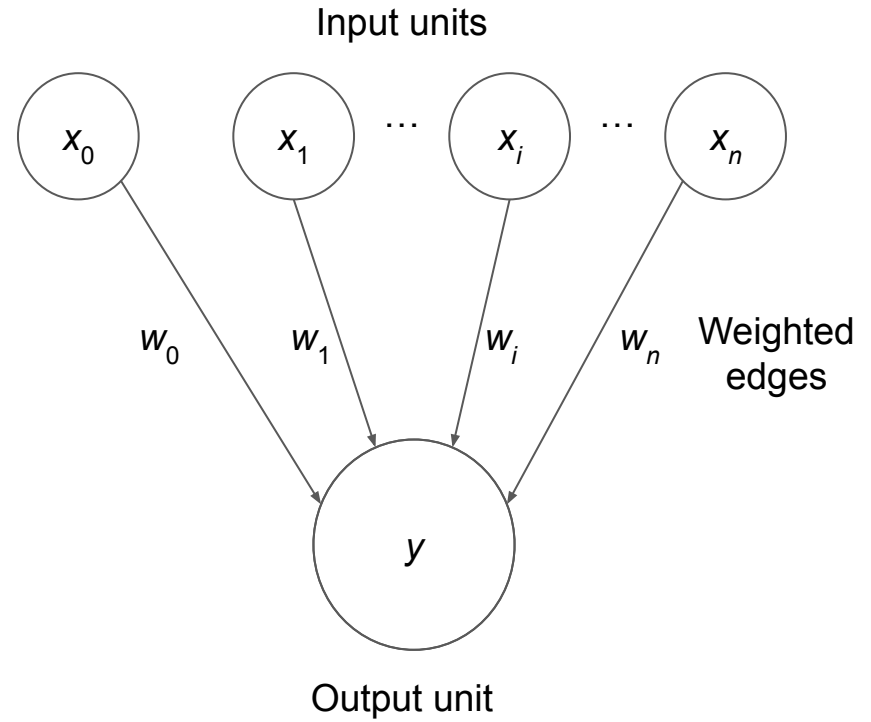


Perceptrons

- Neural networks learn to compute a function using example inputs and outputs
- Single-layer neural network
- Data is fed into input unit units and propagated through the network to the output unit
- A training algorithm strengthens or weakens the connections between the input and the output by modifying the weights

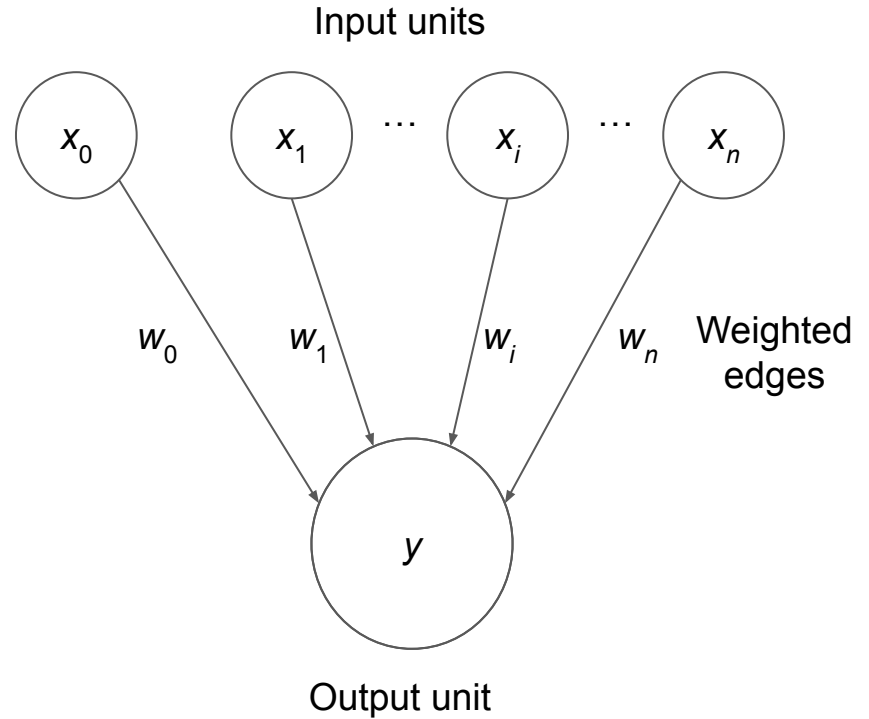


Predicting a Branch



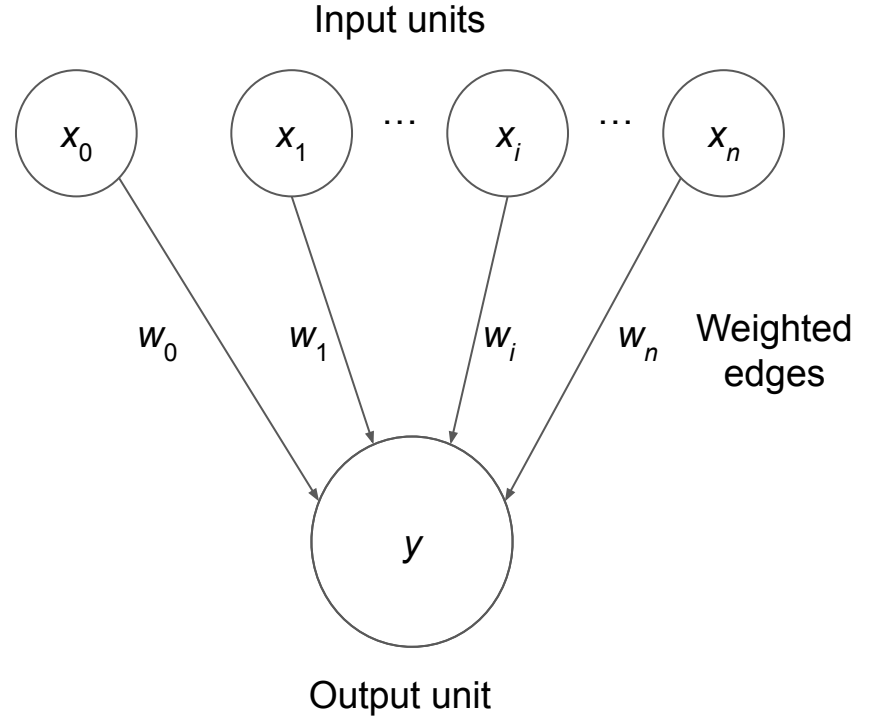
Predicting a Branch

- Input units (x_i)



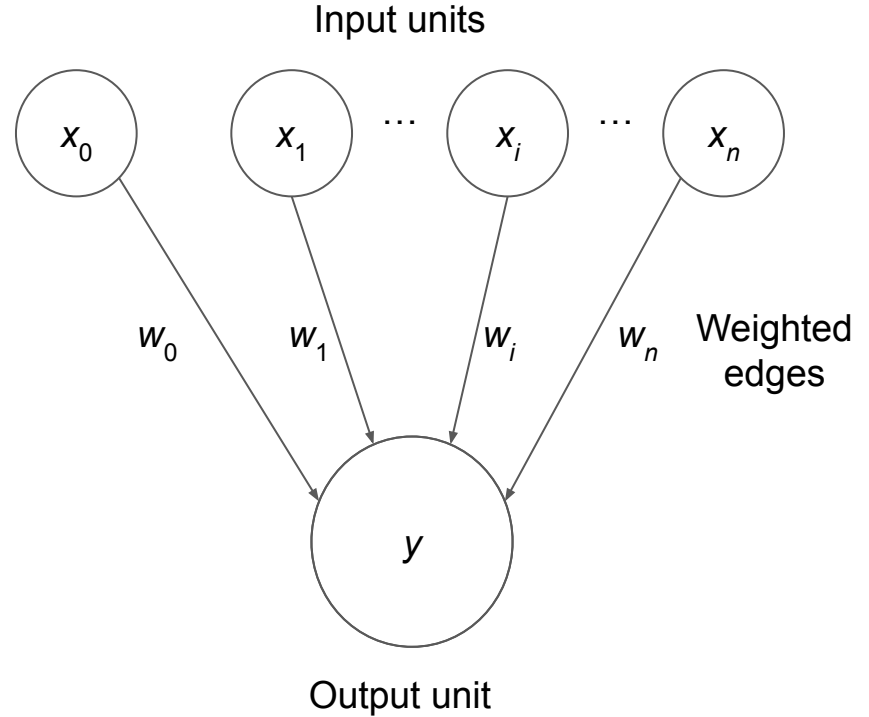
Predicting a Branch

- Input units (x_i)
 - Use bits of global register



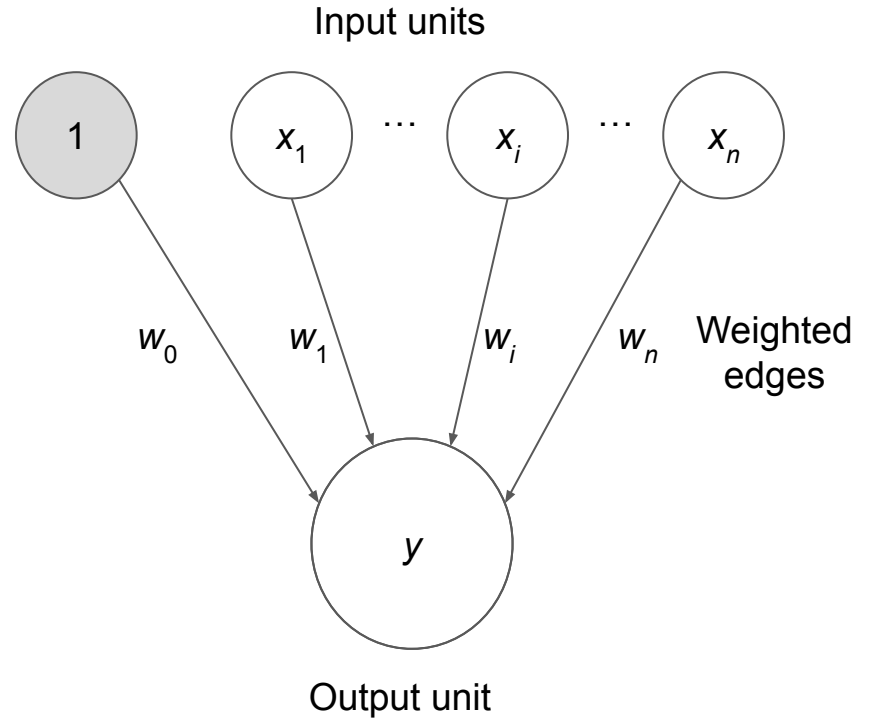
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$



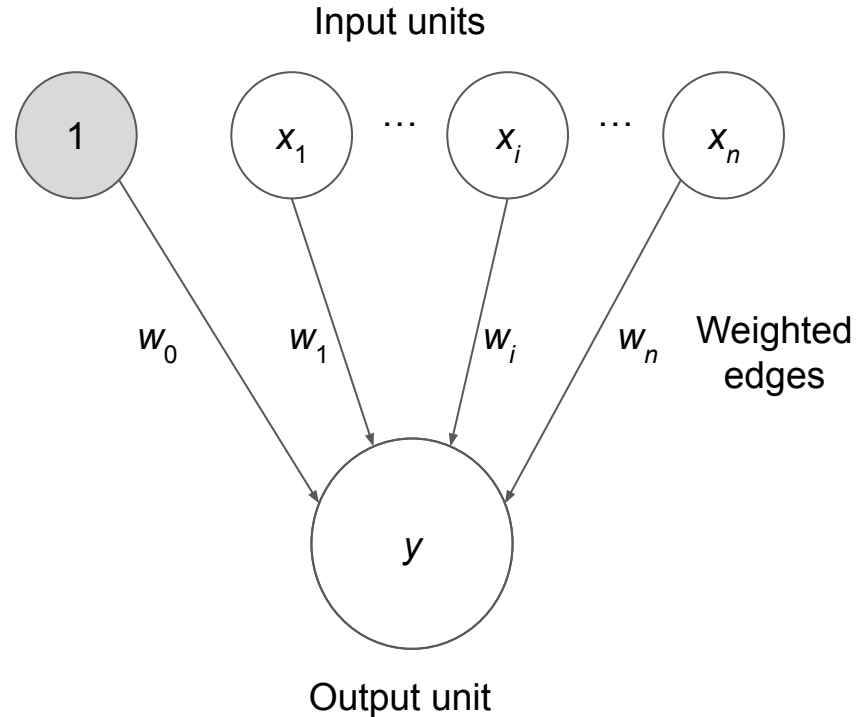
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input



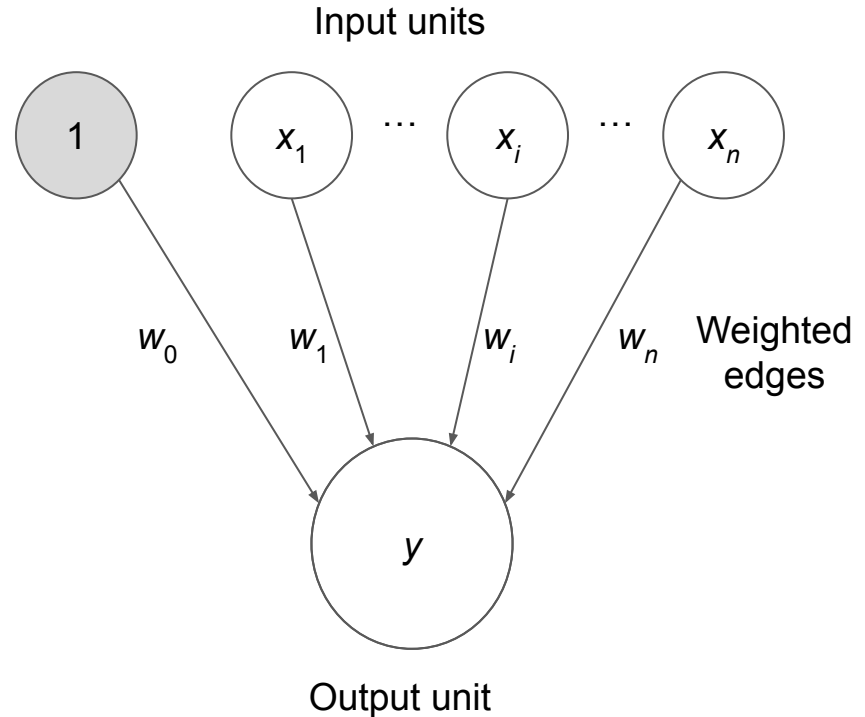
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)



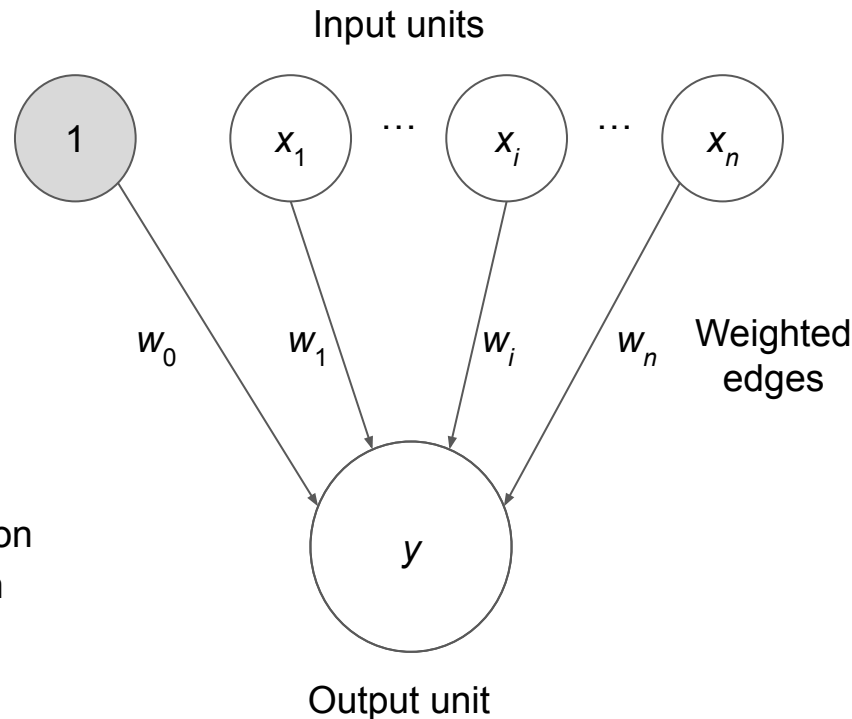
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch



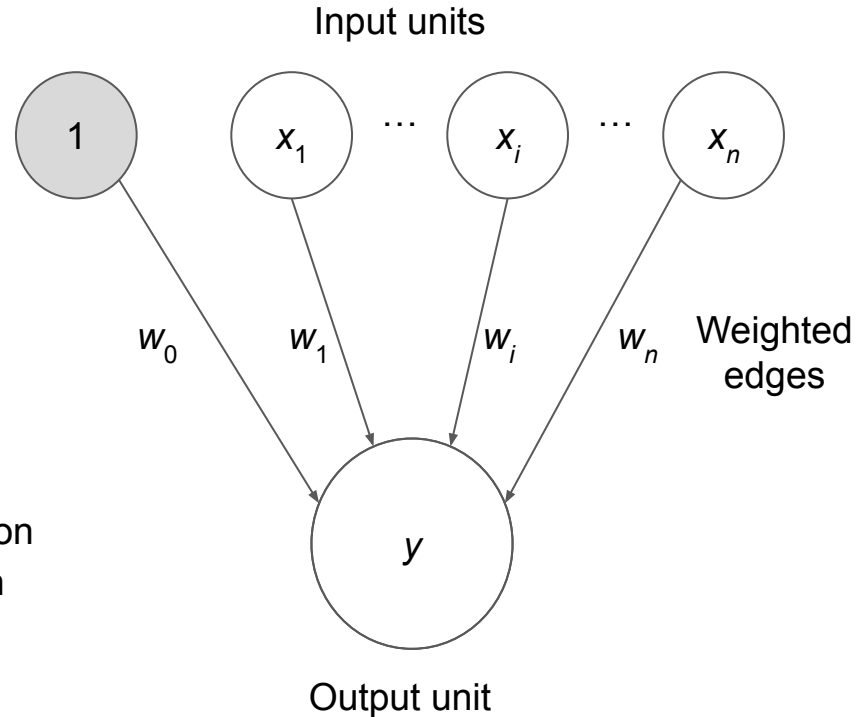
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation



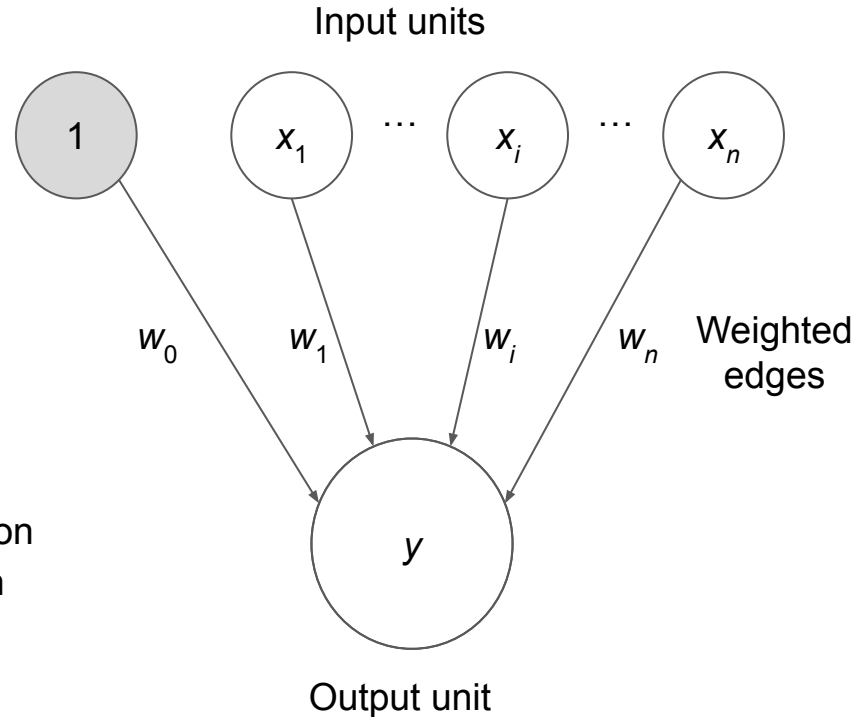
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)



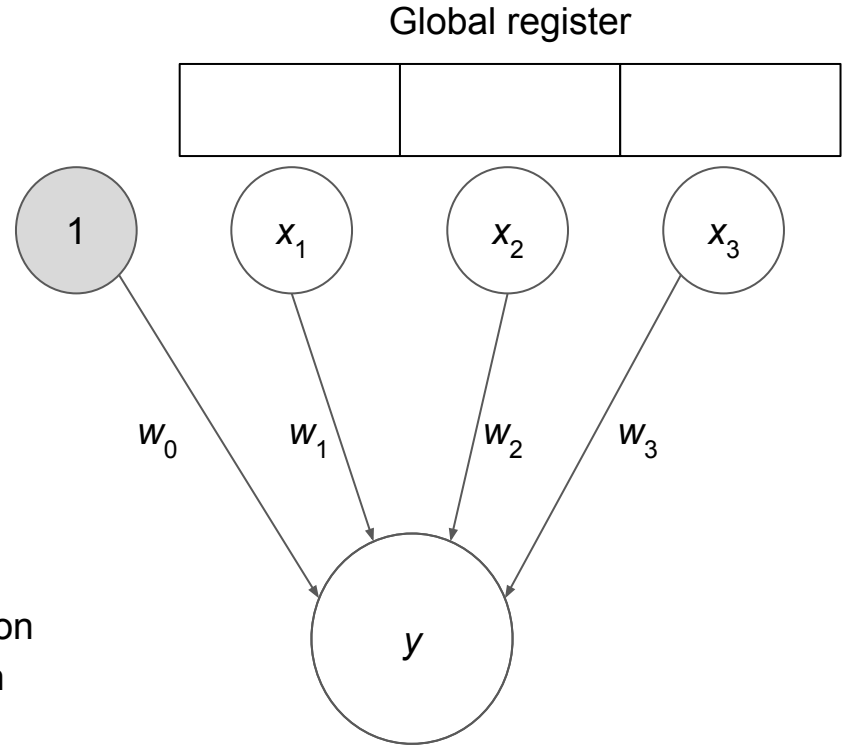
Predicting a Branch

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



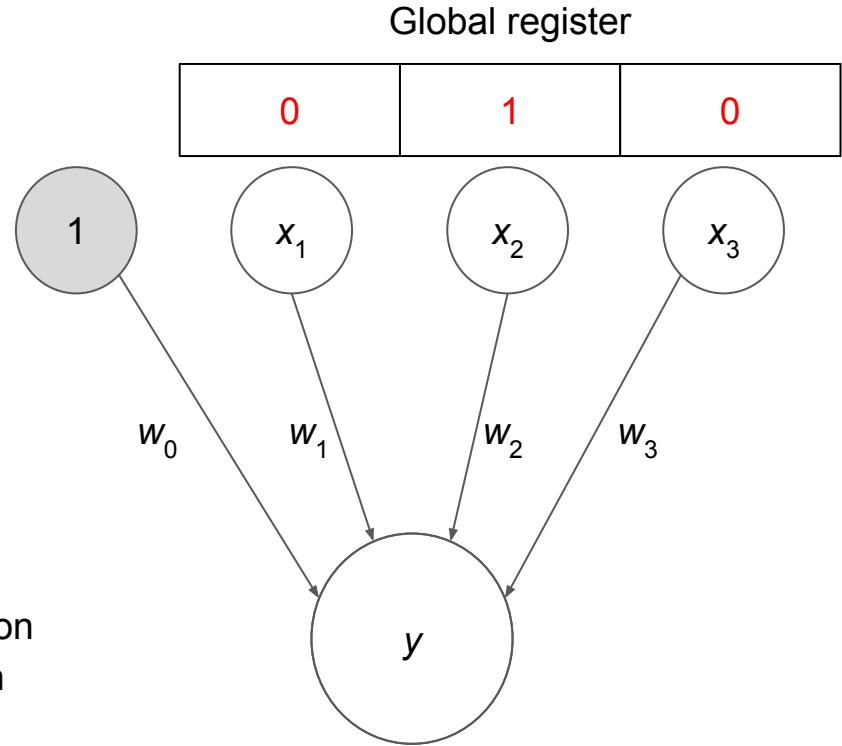
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



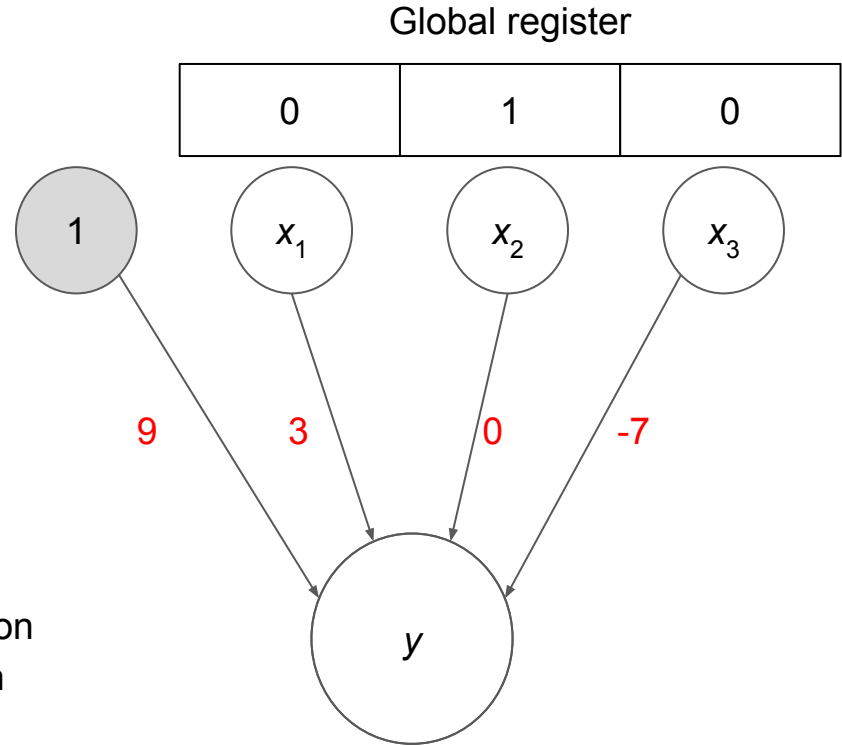
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



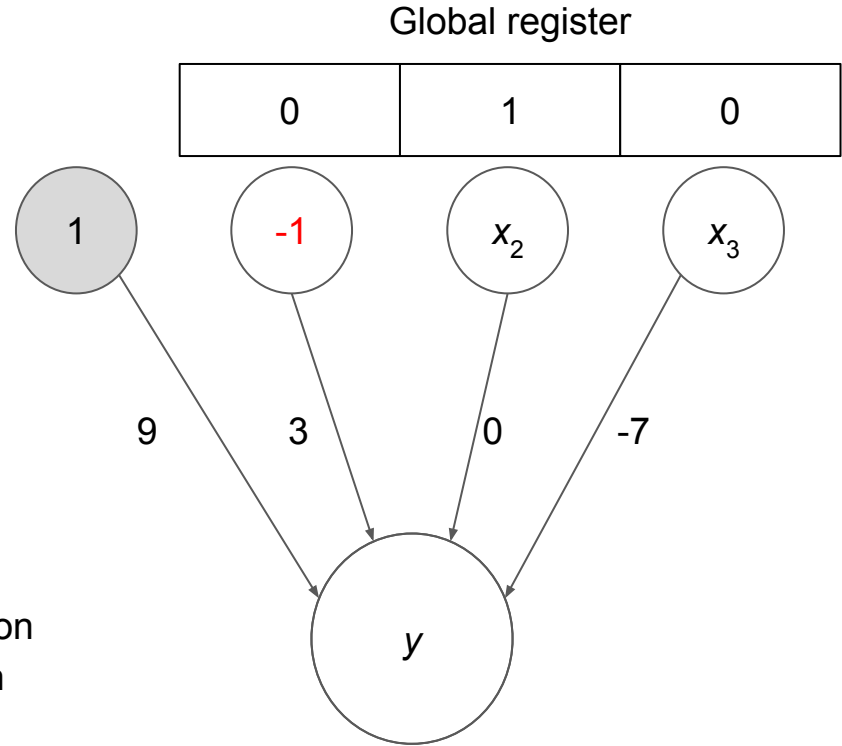
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



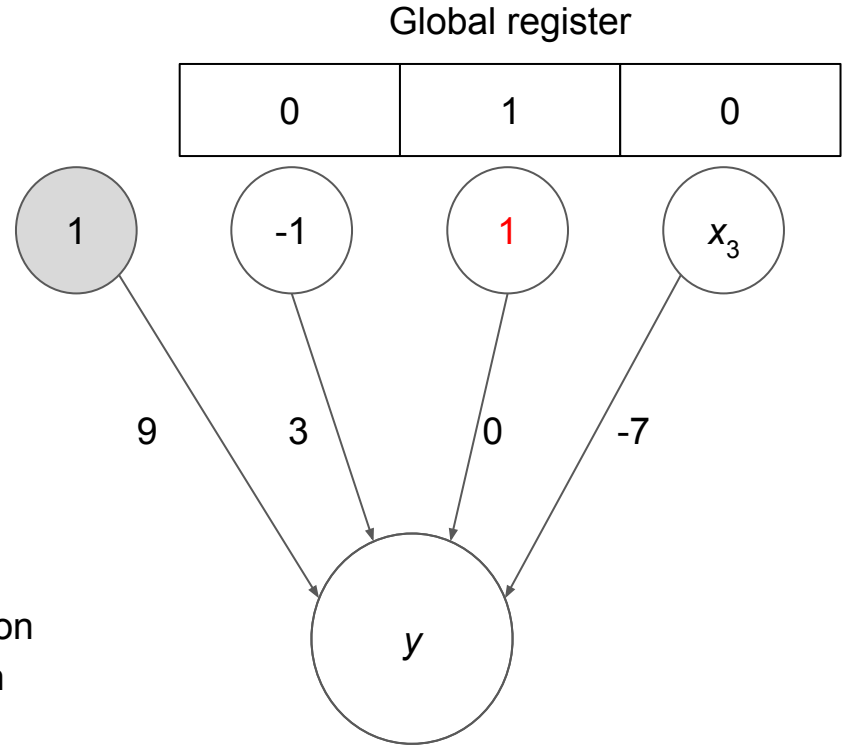
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



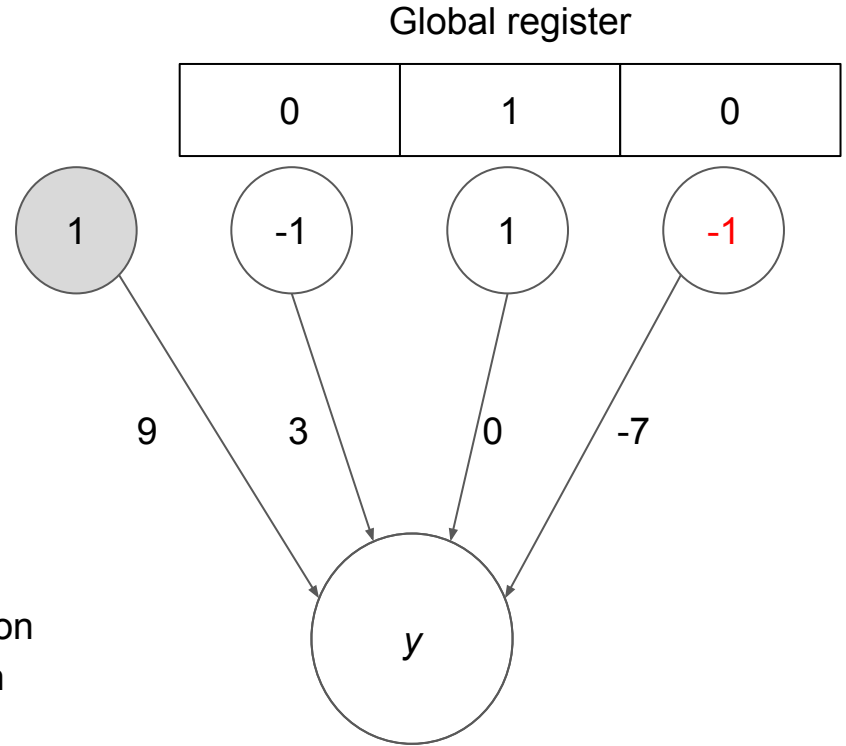
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



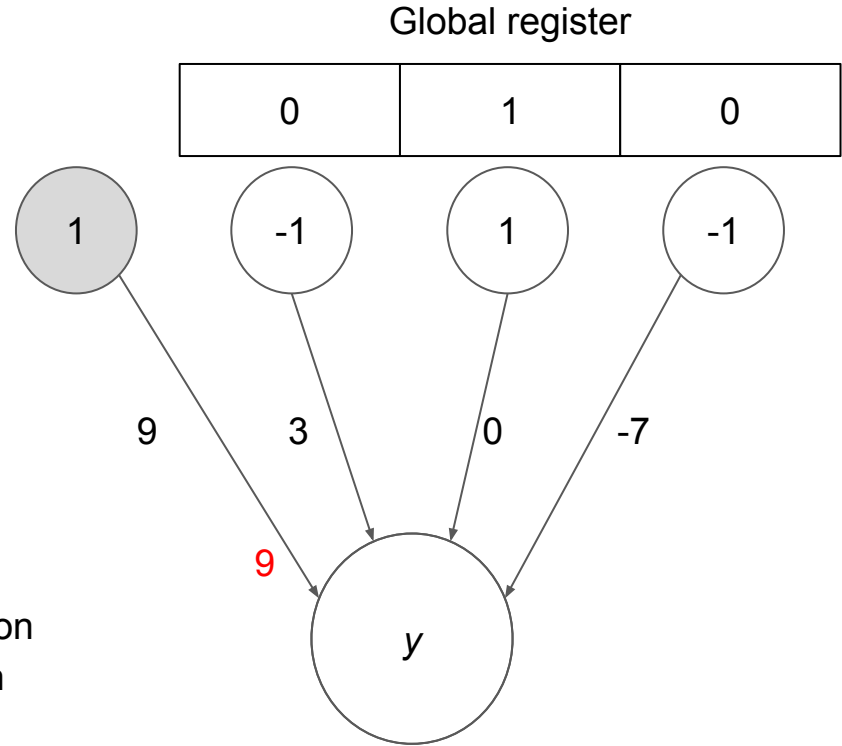
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



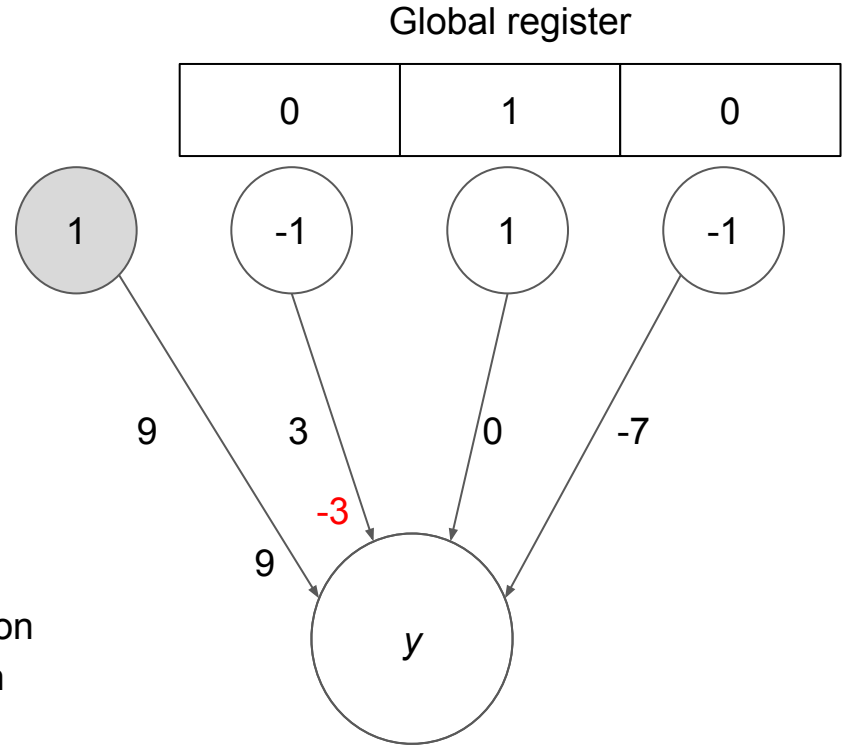
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



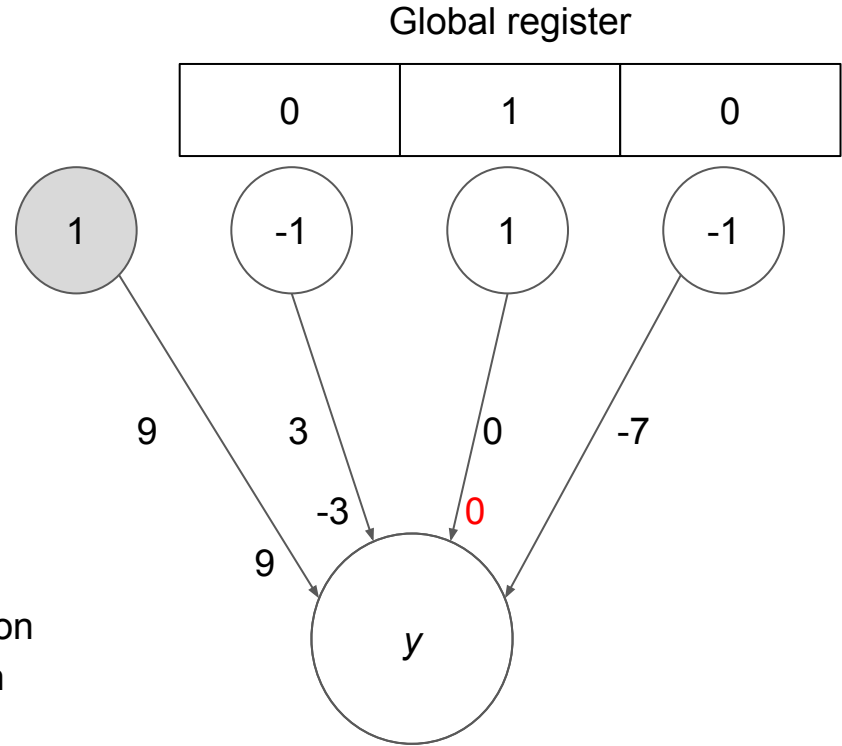
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



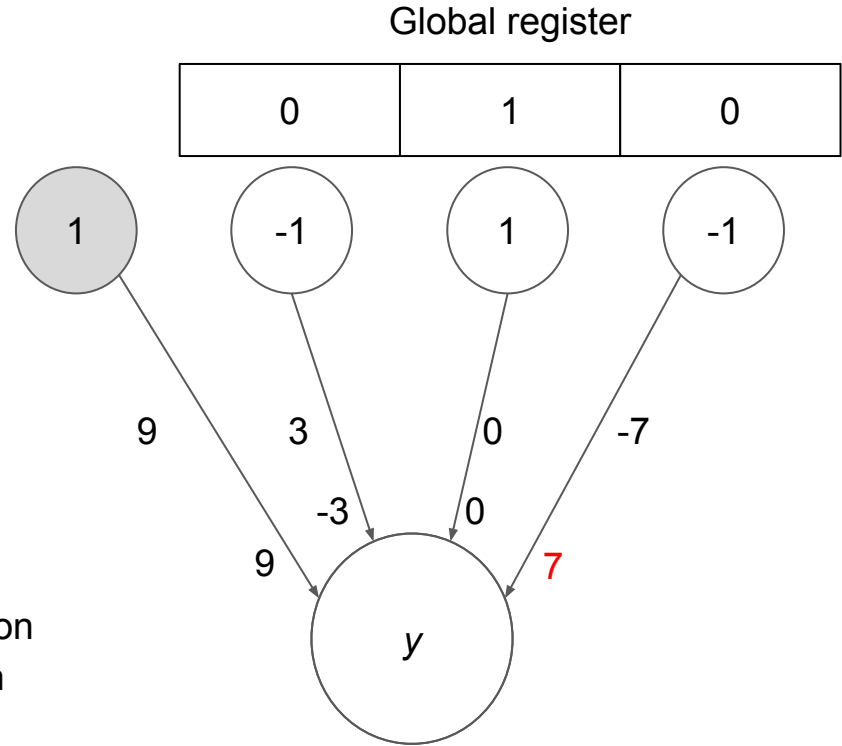
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



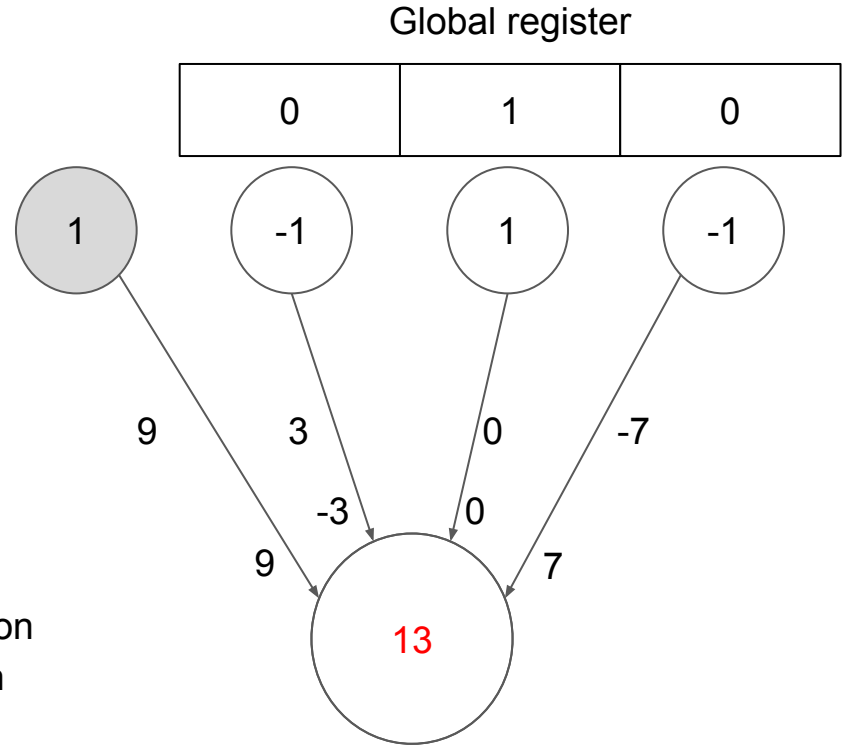
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



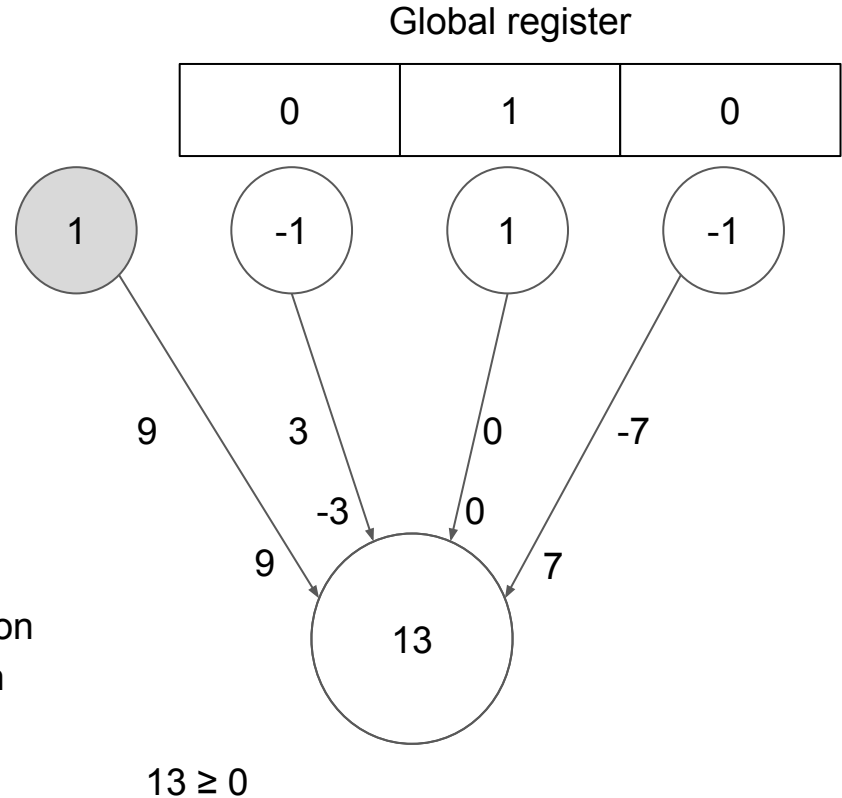
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



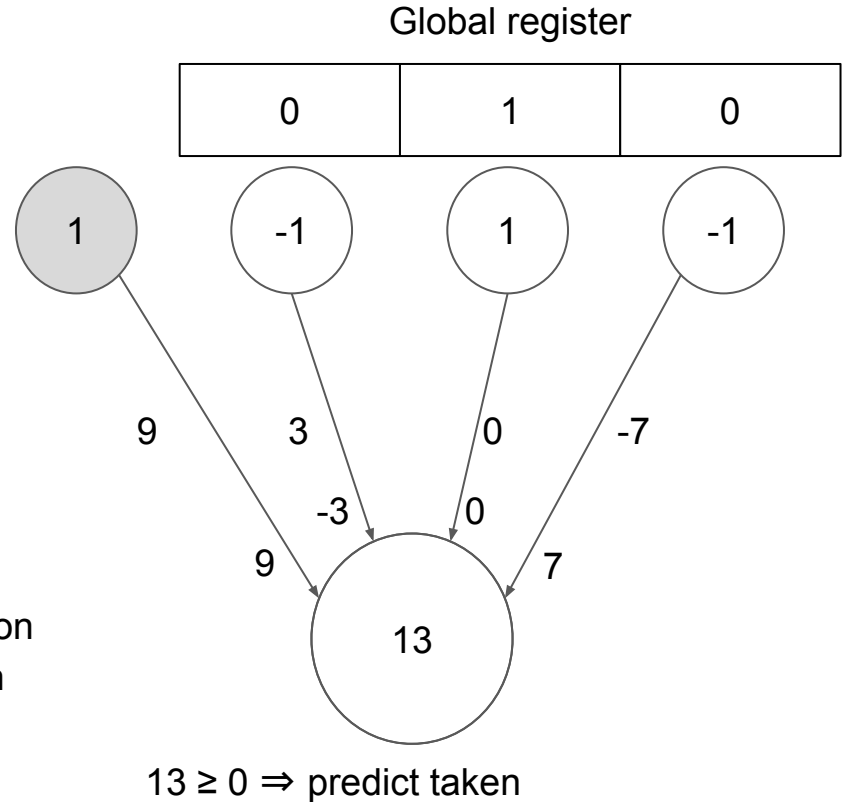
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



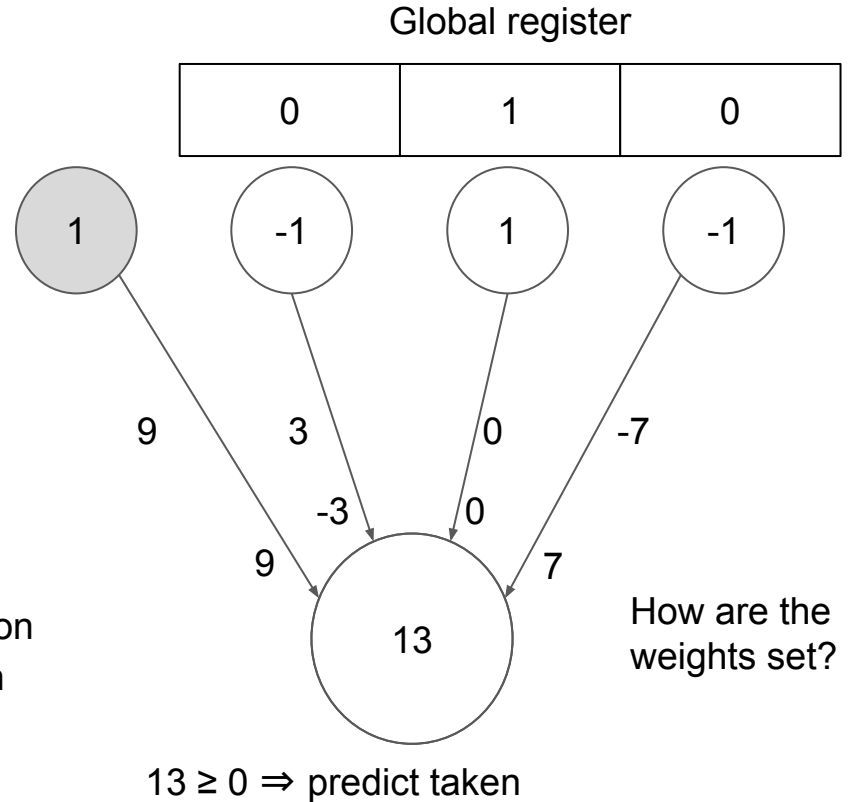
Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken

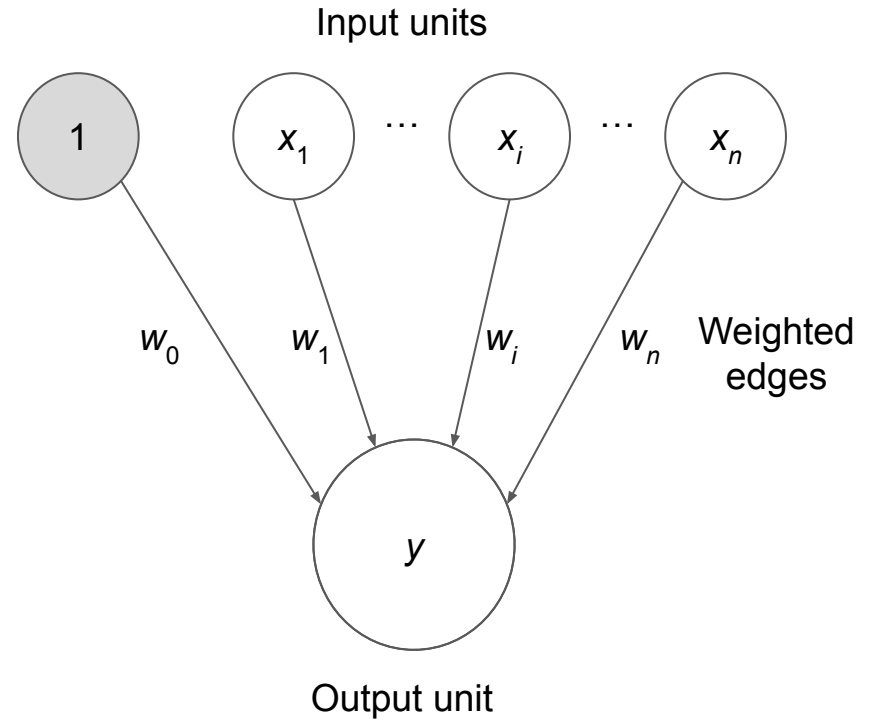


Example

- Input units (x_i)
 - Use bits of global register
 - taken: $x_i = 1$
 - not taken: $x_i = -1$
 - x_0 is always set to 1 as a bias input
- Weights (w_i)
 - Signed integers that represent the correlation between the current branch and the i -th bit of the global register or the bias of the branch
 - very negative: strong negative correlation
 - very positive: strong positive correlation
 - close to 0: weak correlation
- Output (y)
 - $y < 0$: predict not taken
 - $y \geq 0$: predict taken



Training

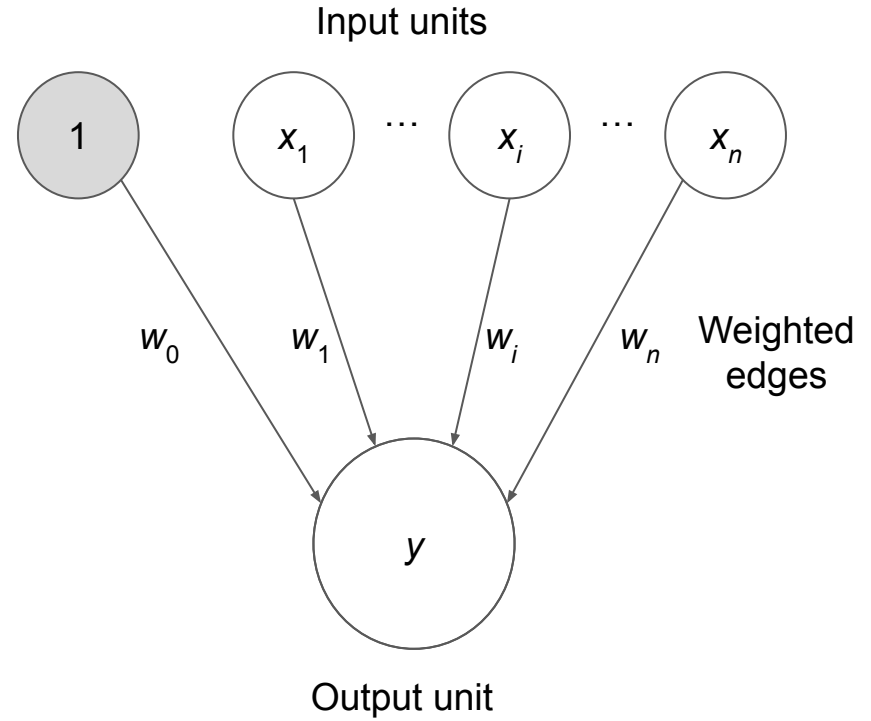


Training

if

then

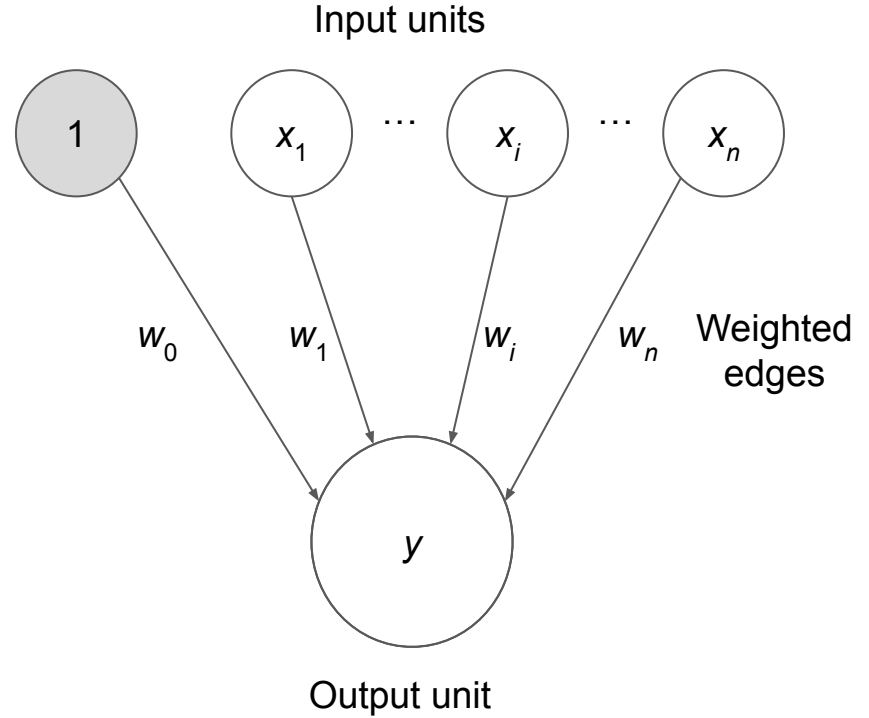
end if



Training

t is 1 if the branch was taken and
-1 otherwise

```
if  $\text{sign}(y) \neq t$  then  
  
end if
```



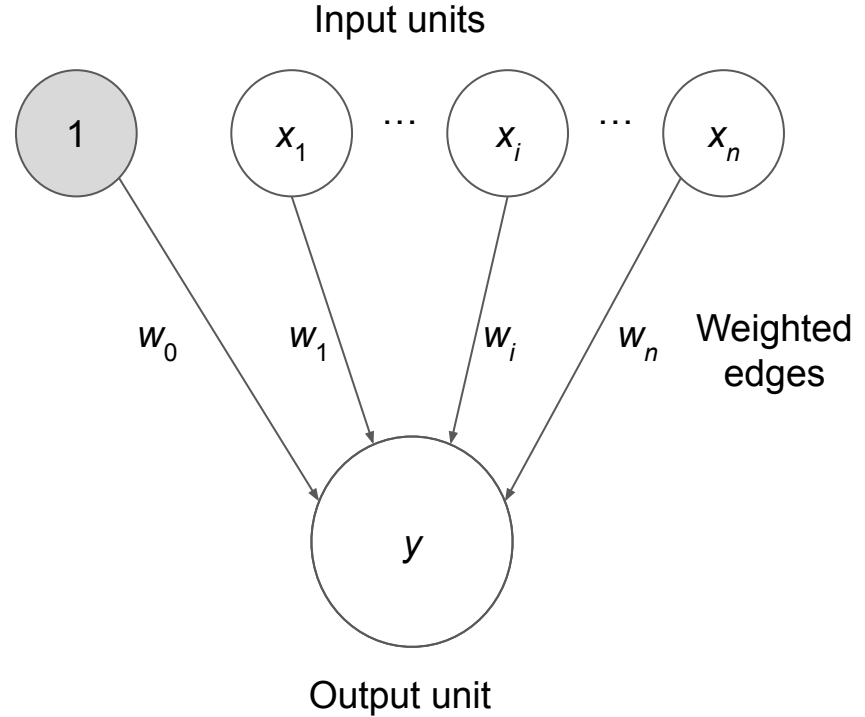
Training

t is 1 if the branch was taken and
-1 otherwise

θ is the threshold used to decide when
enough training has been done

if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**

end if

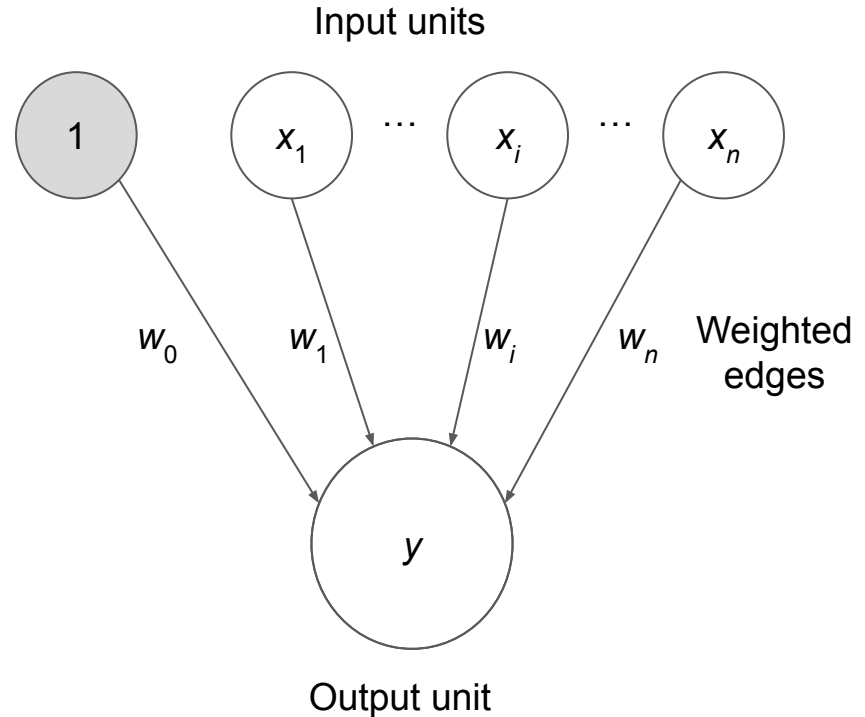


Training

t is 1 if the branch was taken and
-1 otherwise

θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
  
    end for  
end if
```

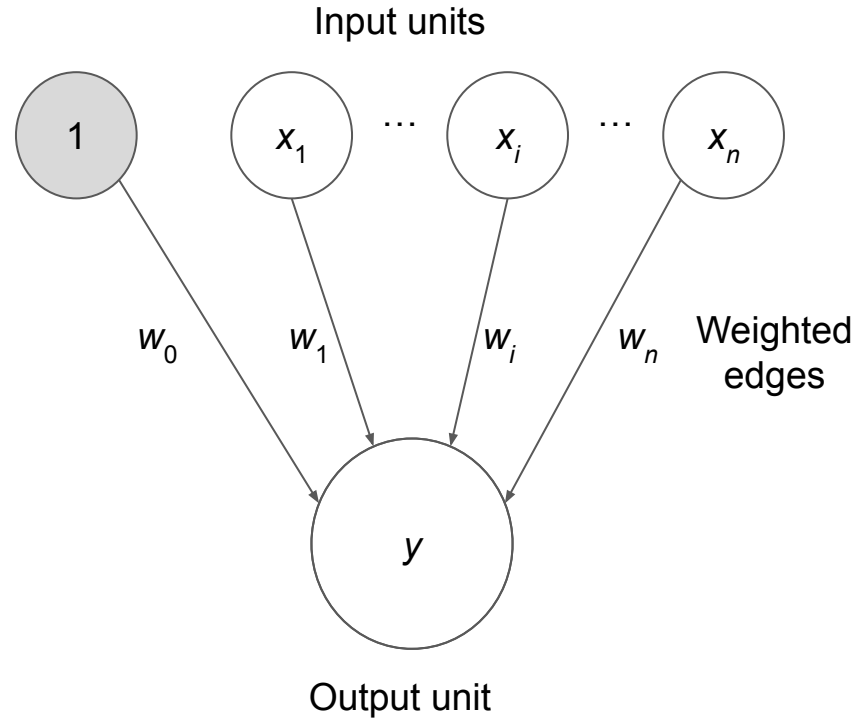


Training

t is 1 if the branch was taken and
-1 otherwise

θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```

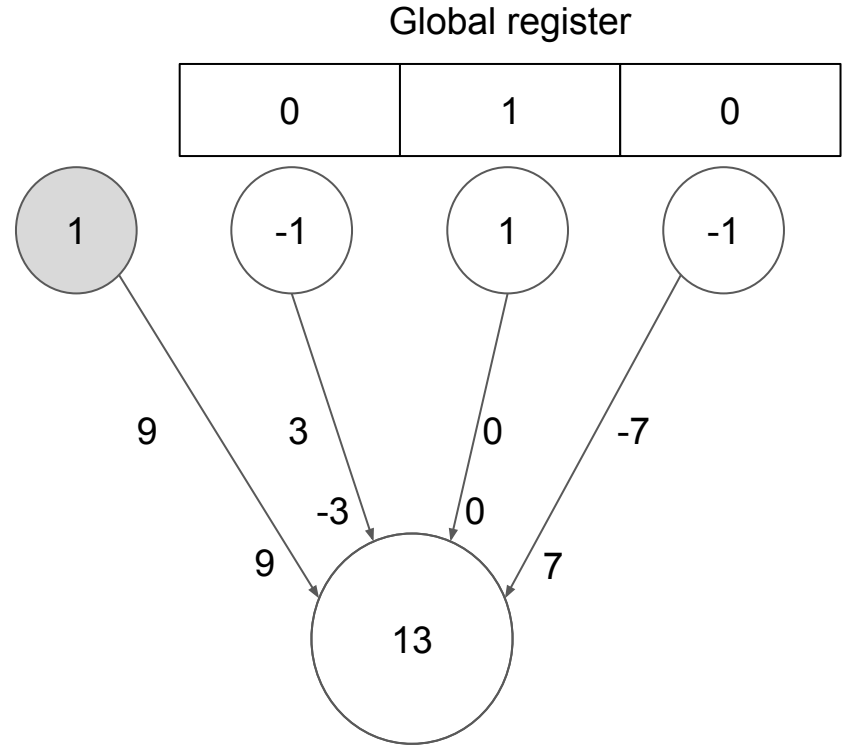


Example

t is 1 if the branch was taken and
-1 otherwise

θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```

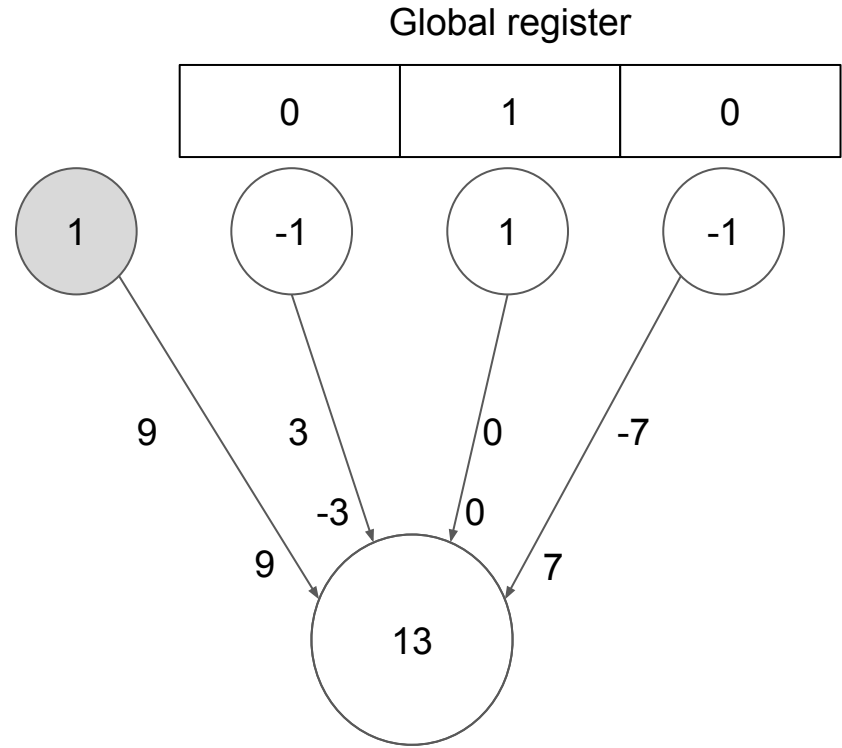


Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



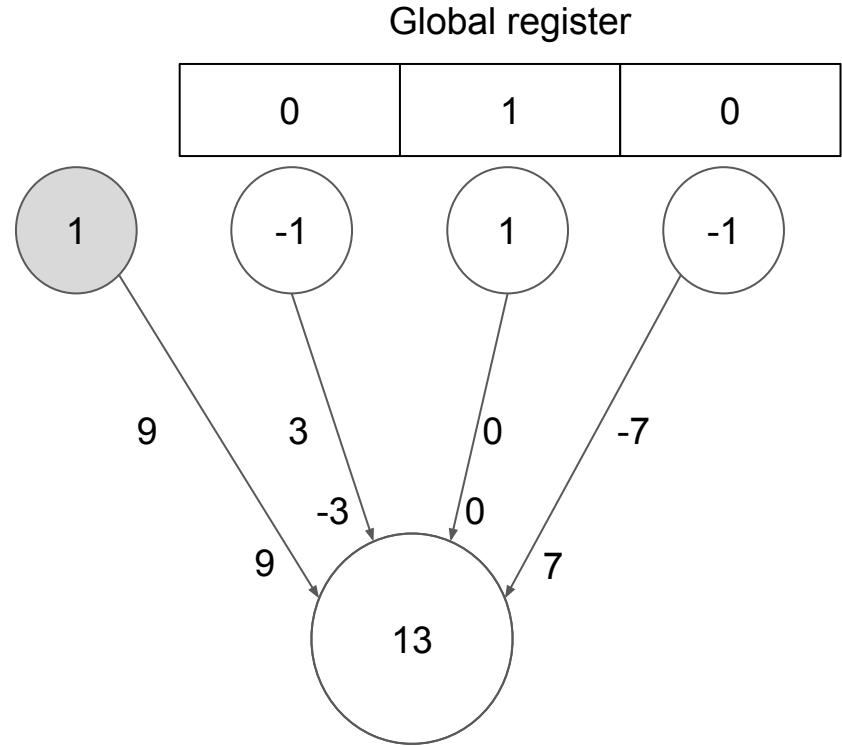
Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



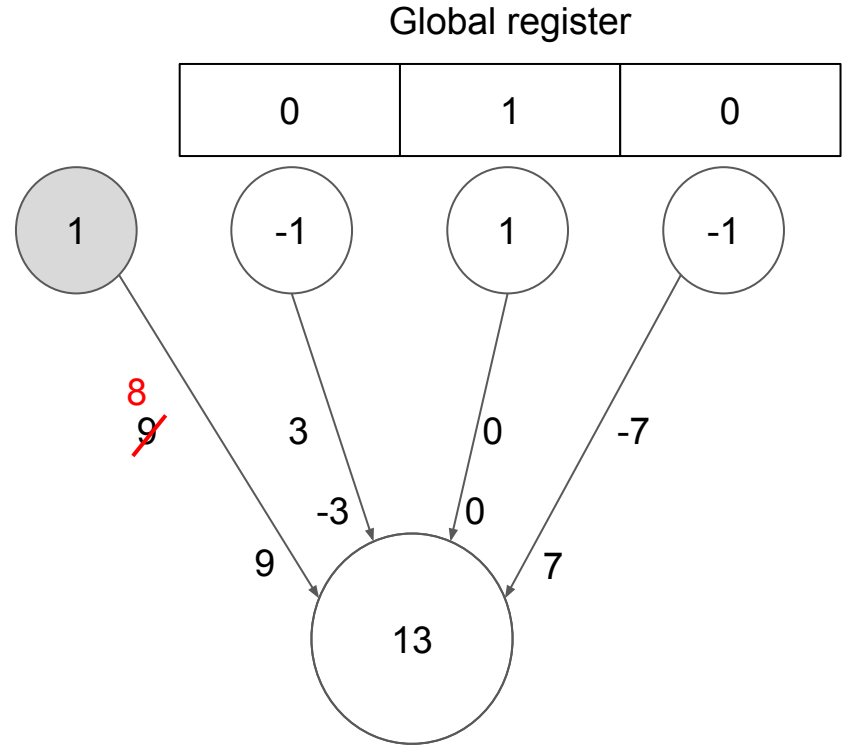
Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**
 for $i := 0$ **to** n **do**
 $w_i := w_i + tx_i$
 end for
end if



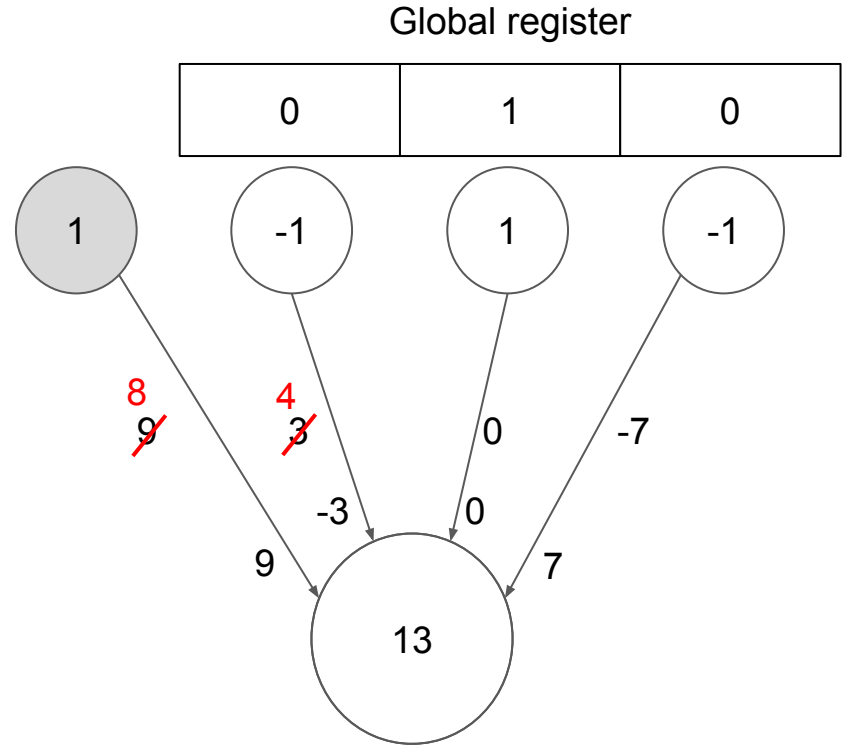
Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**
 for $i := 0$ **to** n **do**
 $w_i := w_i + tx_i$
 end for
end if



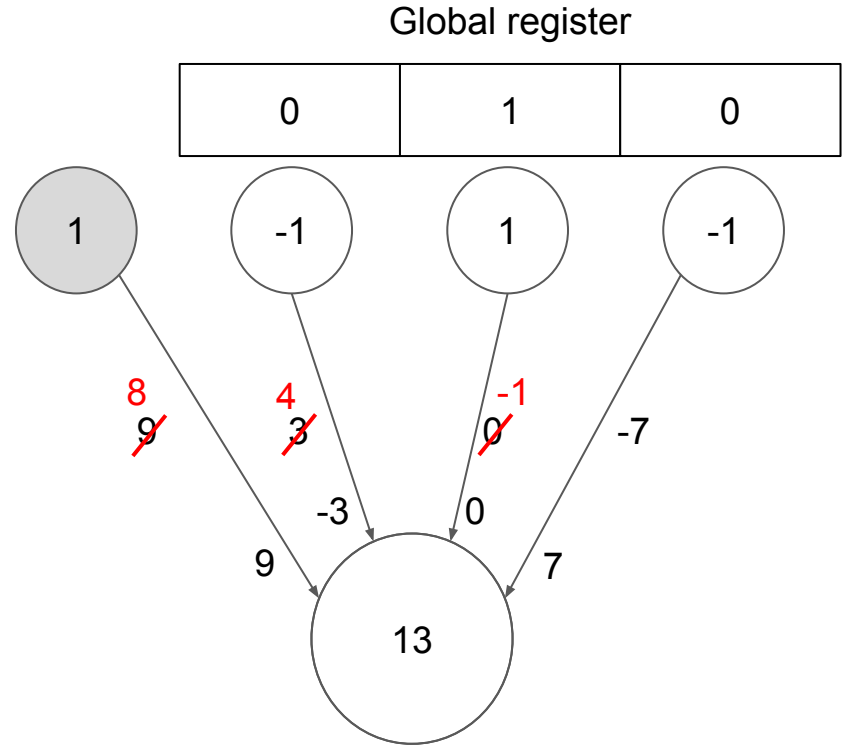
Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



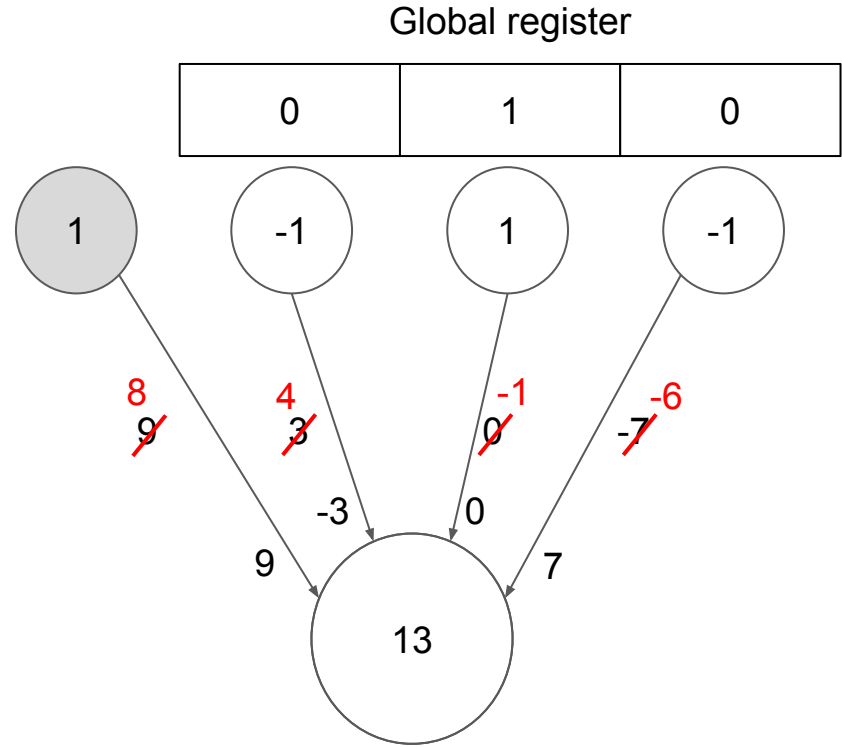
Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

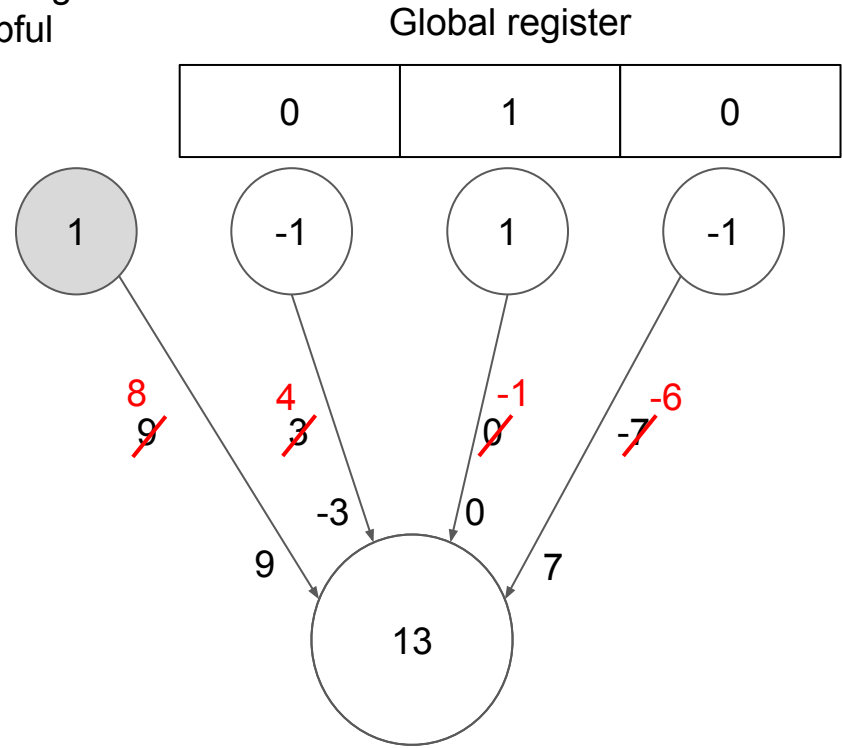
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



Let $\Theta = 15$ and $t = -1$ (the branch was not taken)

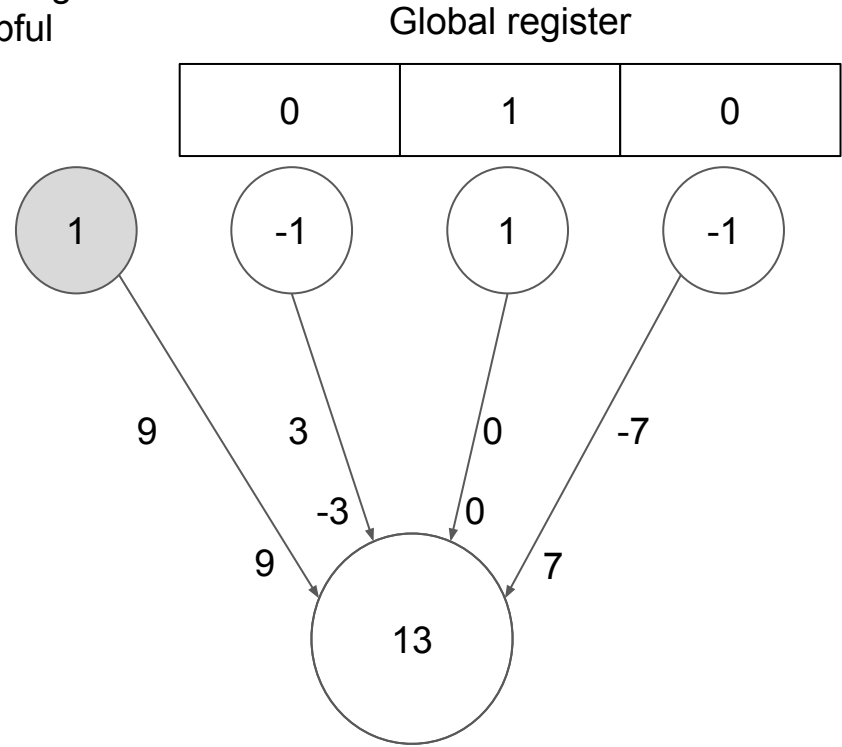
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



Let $\Theta = 15$ and $t = 1$ (the branch **was taken**)

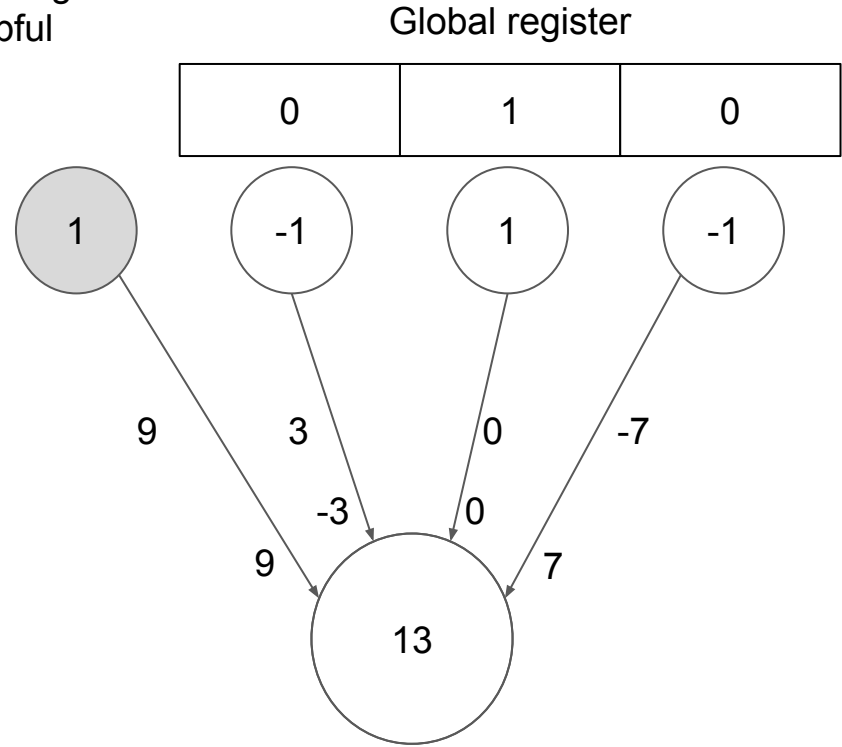
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



Let $\Theta = 15$ and $t = 1$ (the branch **was taken**)

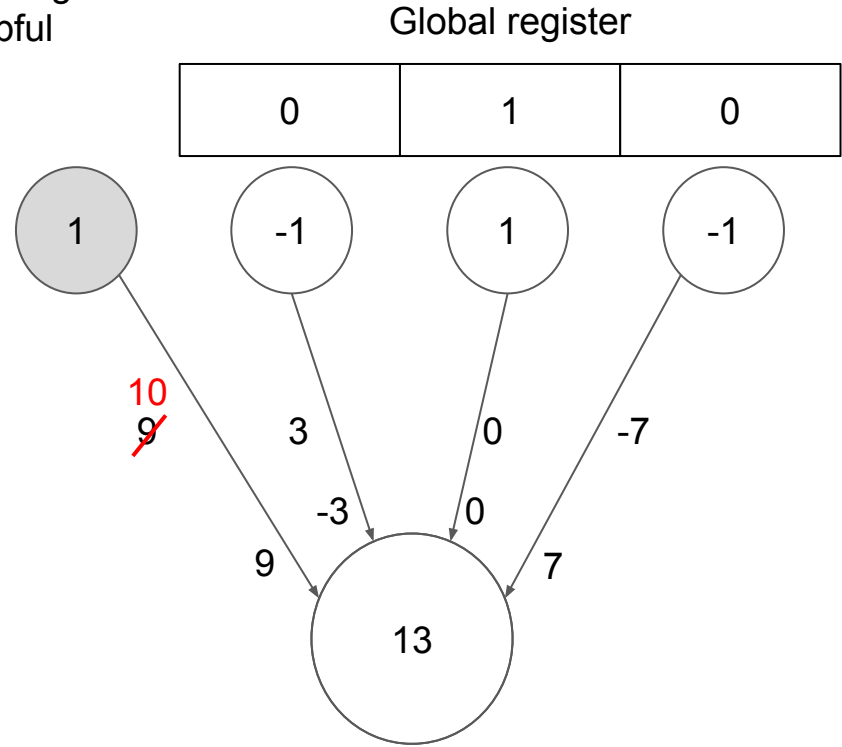
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then
  for  $i := 0$  to  $n$  do
     $w_i := w_i + tx_i$ 
  end for
end if
```



Let $\Theta = 15$ and $t = 1$ (the branch **was taken**)

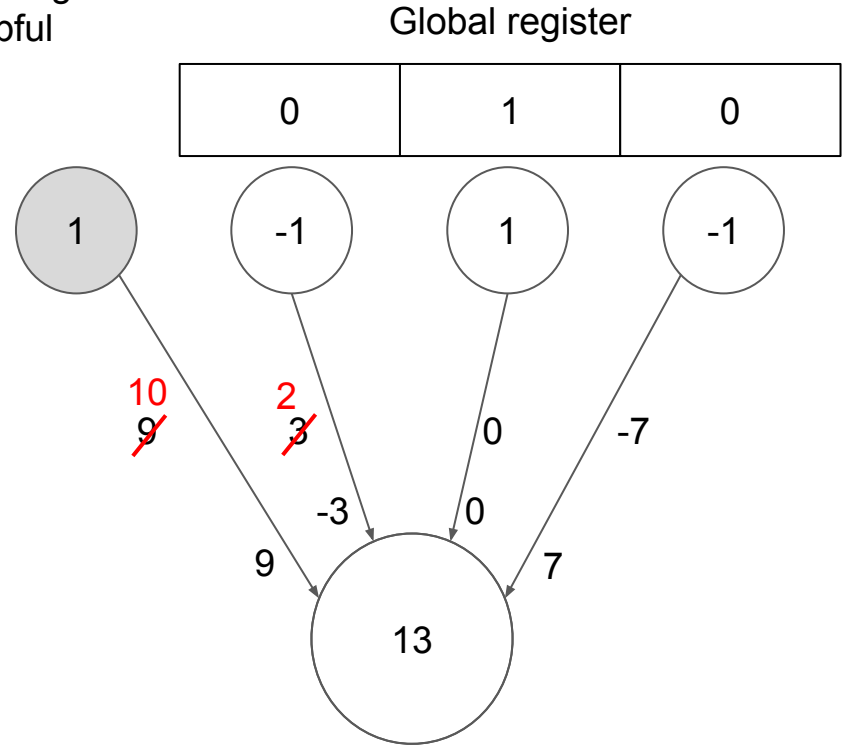
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then
  for  $i := 0$  to  $n$  do
     $w_i := w_i + tx_i$ 
  end for
end if
```



Let $\Theta = 15$ and $t = 1$ (the branch **was taken**)

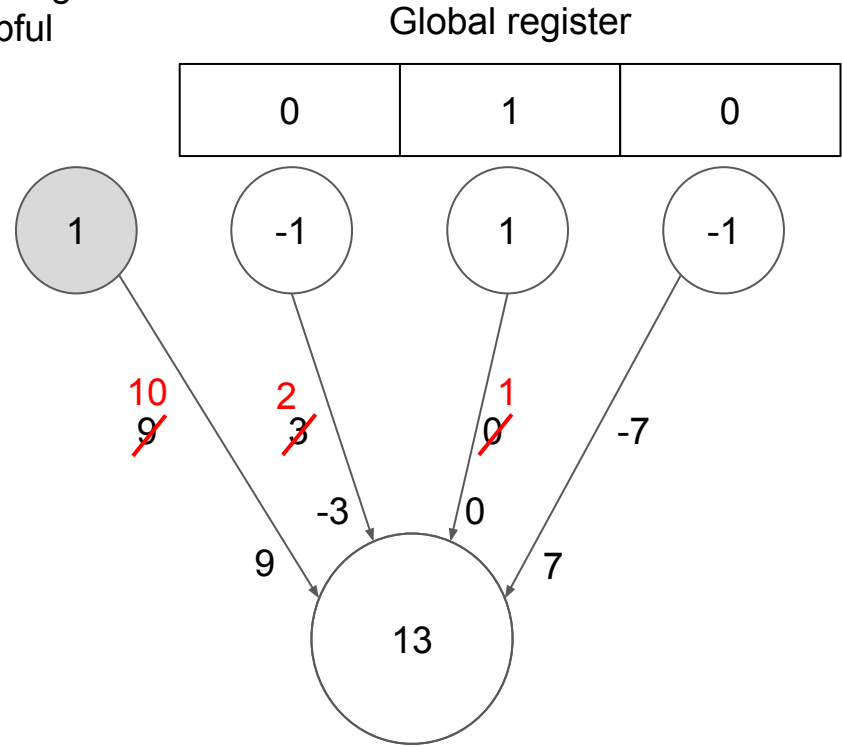
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



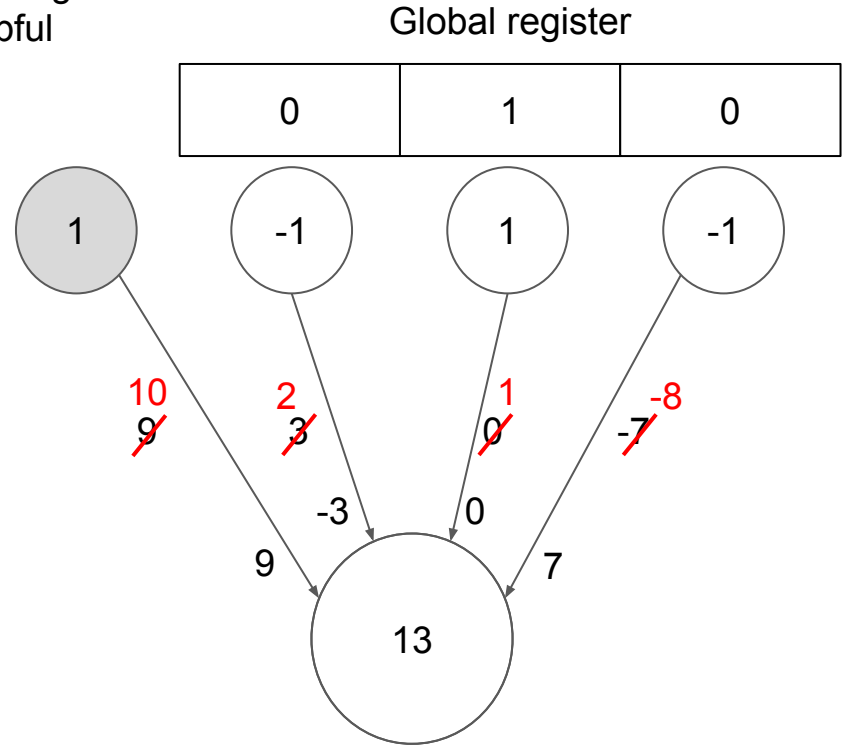
Example

When training, helpful weights get further from zero and unhelpful weights get closer to zero.

t is 1 if the branch was taken and -1 otherwise

Θ is the threshold used to decide when enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```

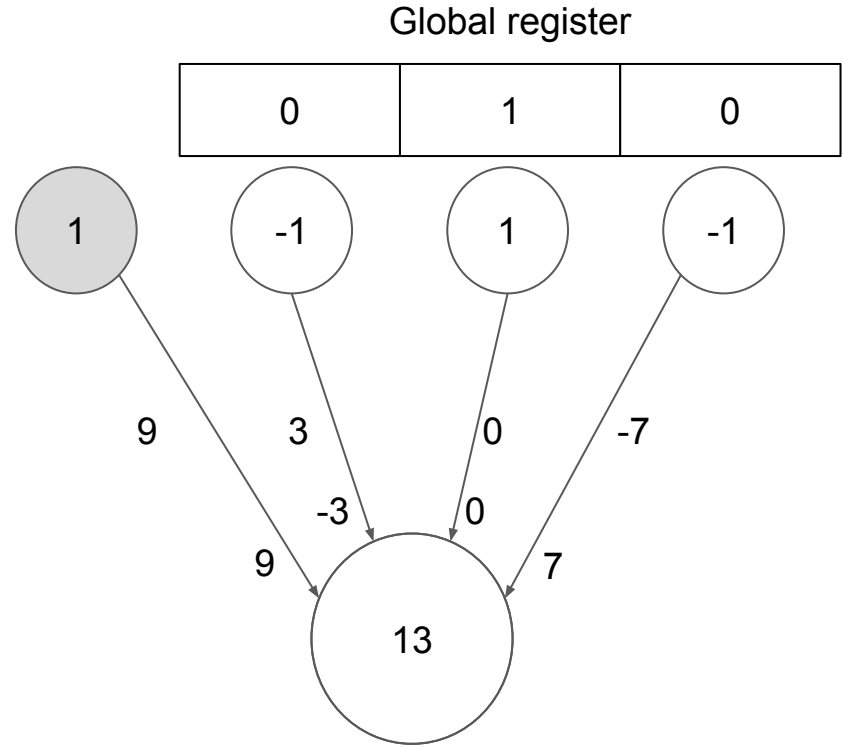


Example

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



Let $\Theta = 10$ and $t = 1$ (the branch was taken)

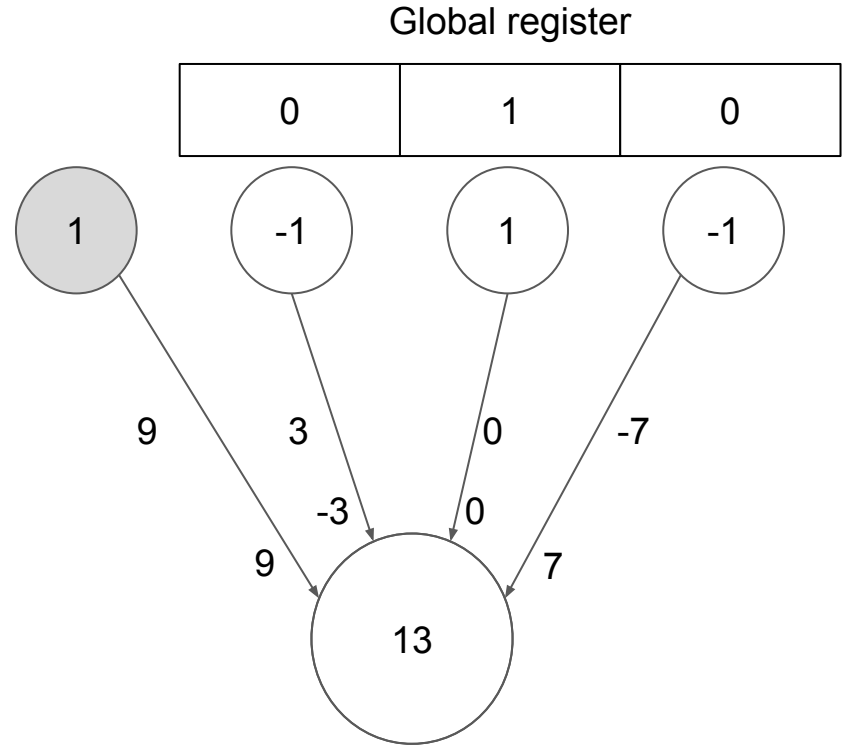
Example

When above the threshold,
only train the weights if the
prediction was **incorrect**

t is 1 if the branch was taken and
-1 otherwise

Θ is the threshold used to decide when
enough training has been done

```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



Let $\Theta = 10$ and $t = 1$ (the branch was taken)

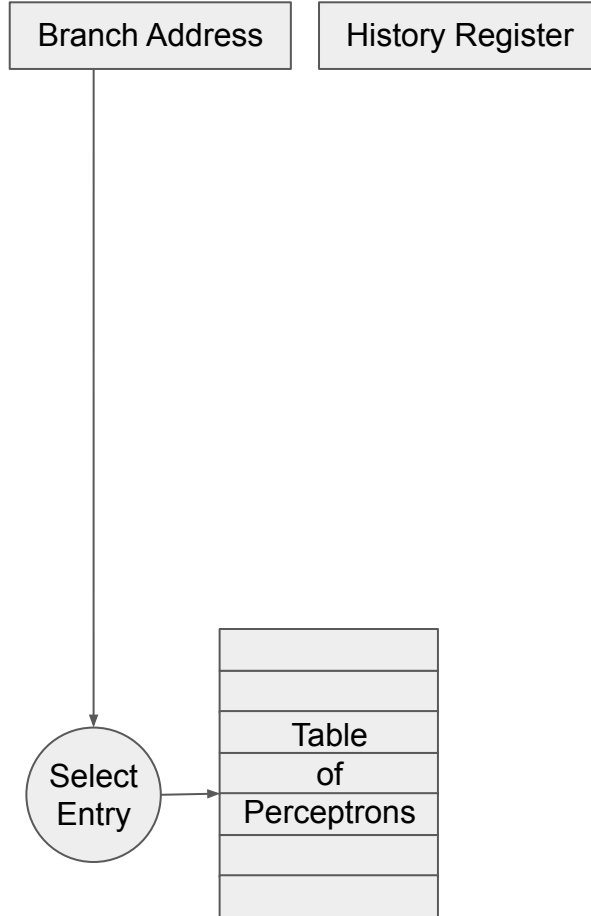
History Register

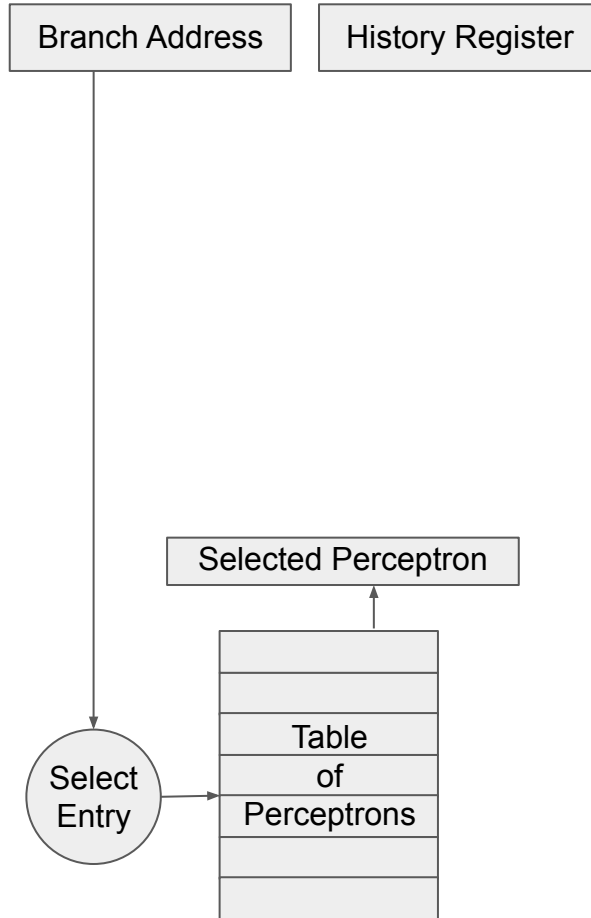
Table
of
Perceptrons

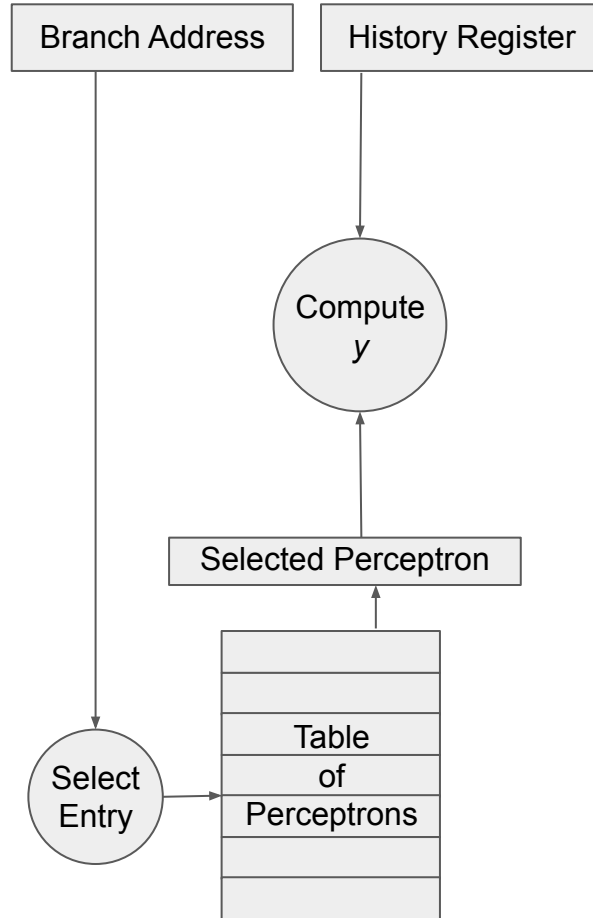
Branch Address

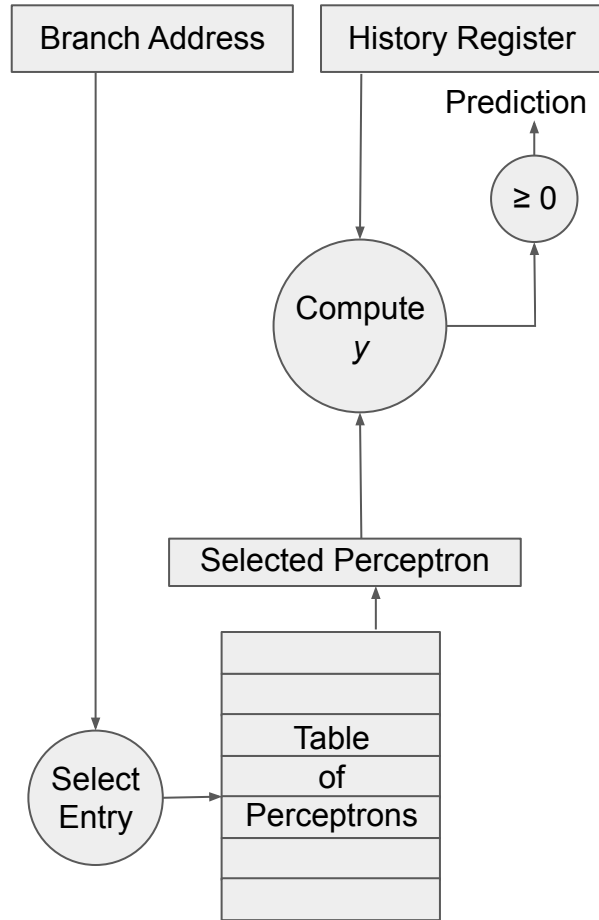
History Register

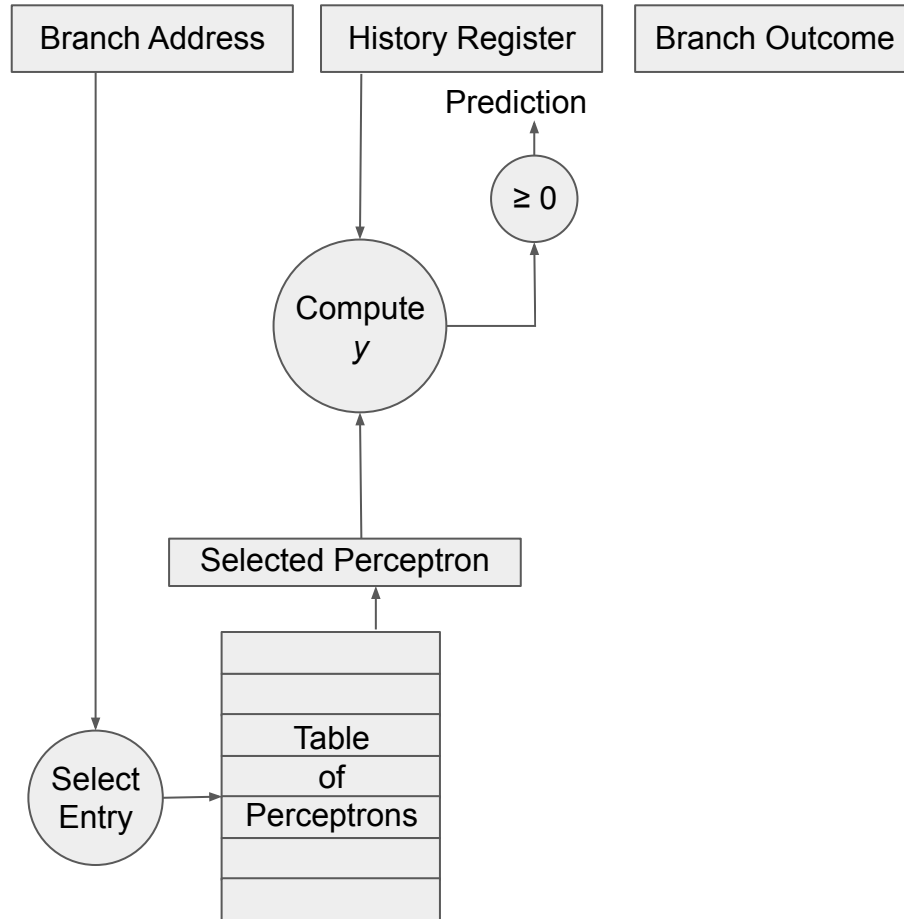
Table
of
Perceptrons

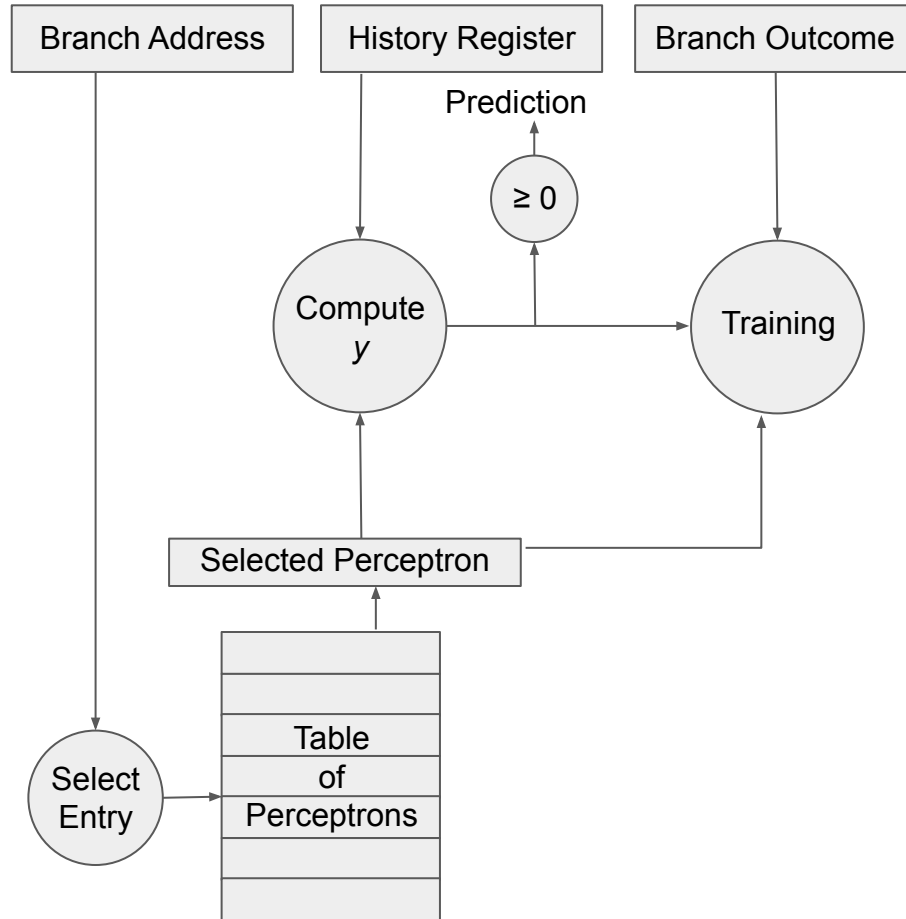


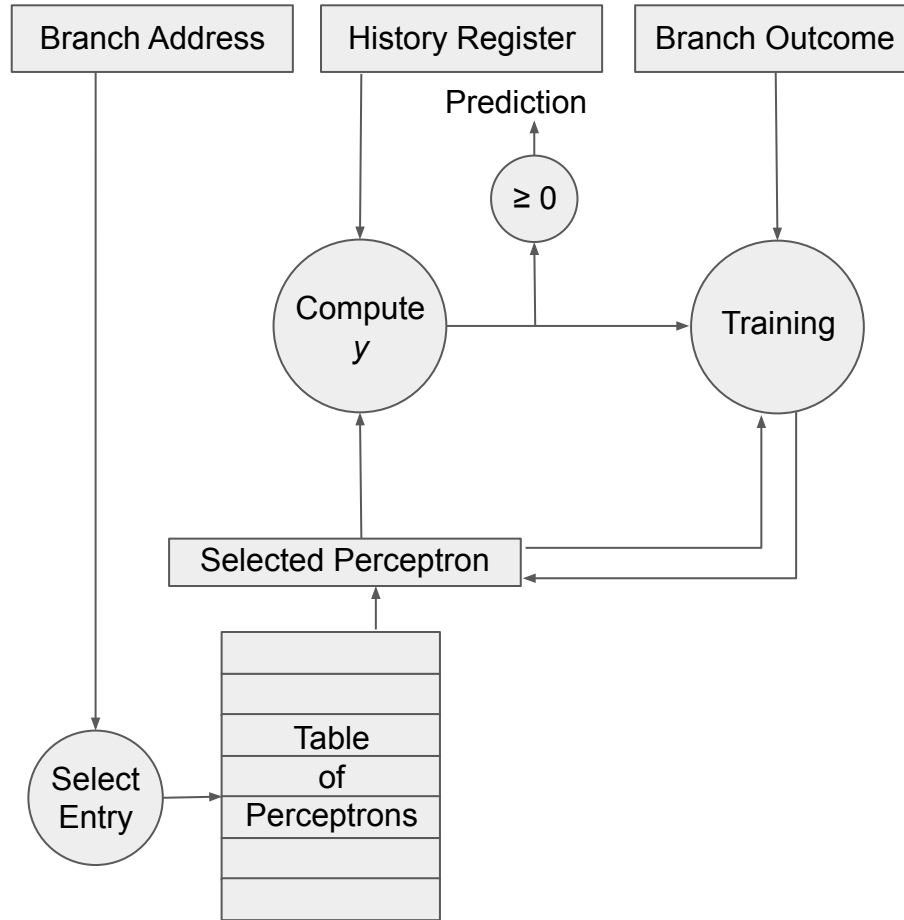


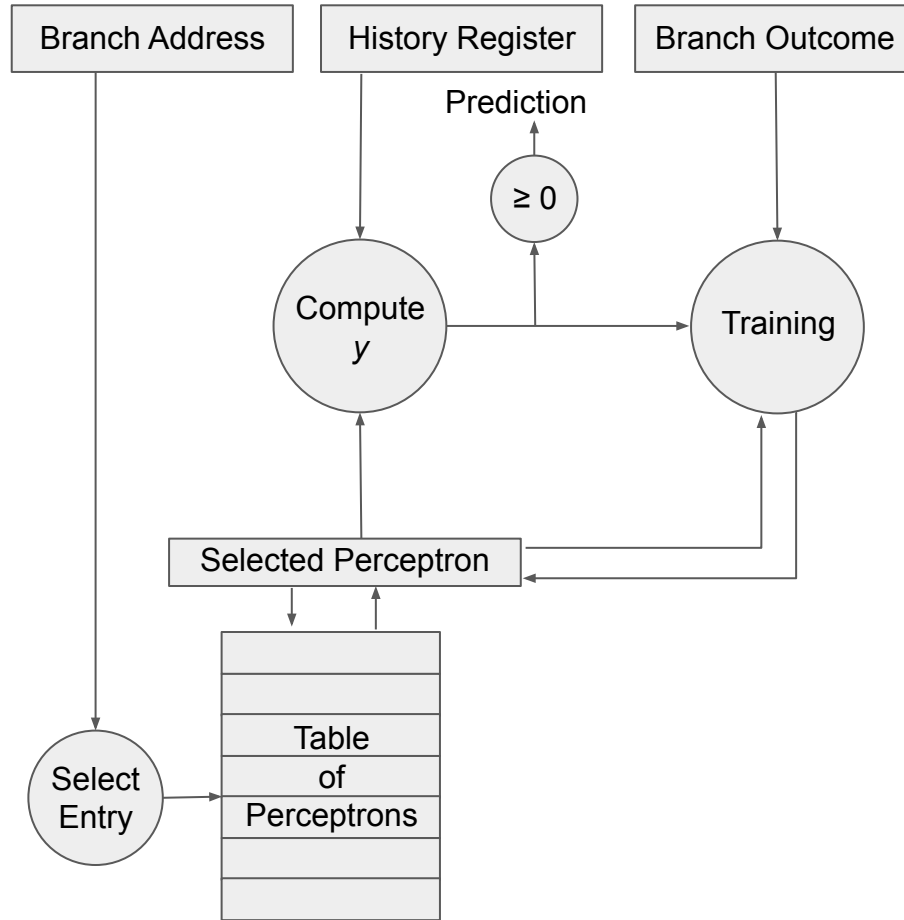


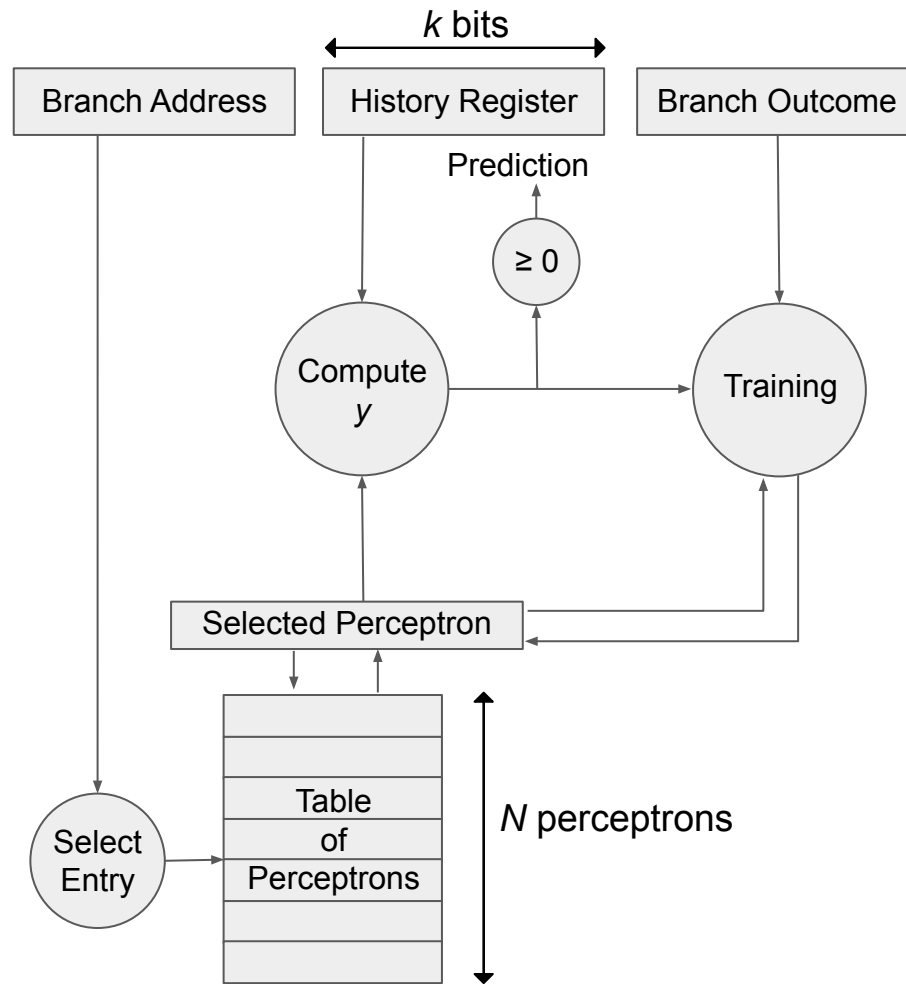


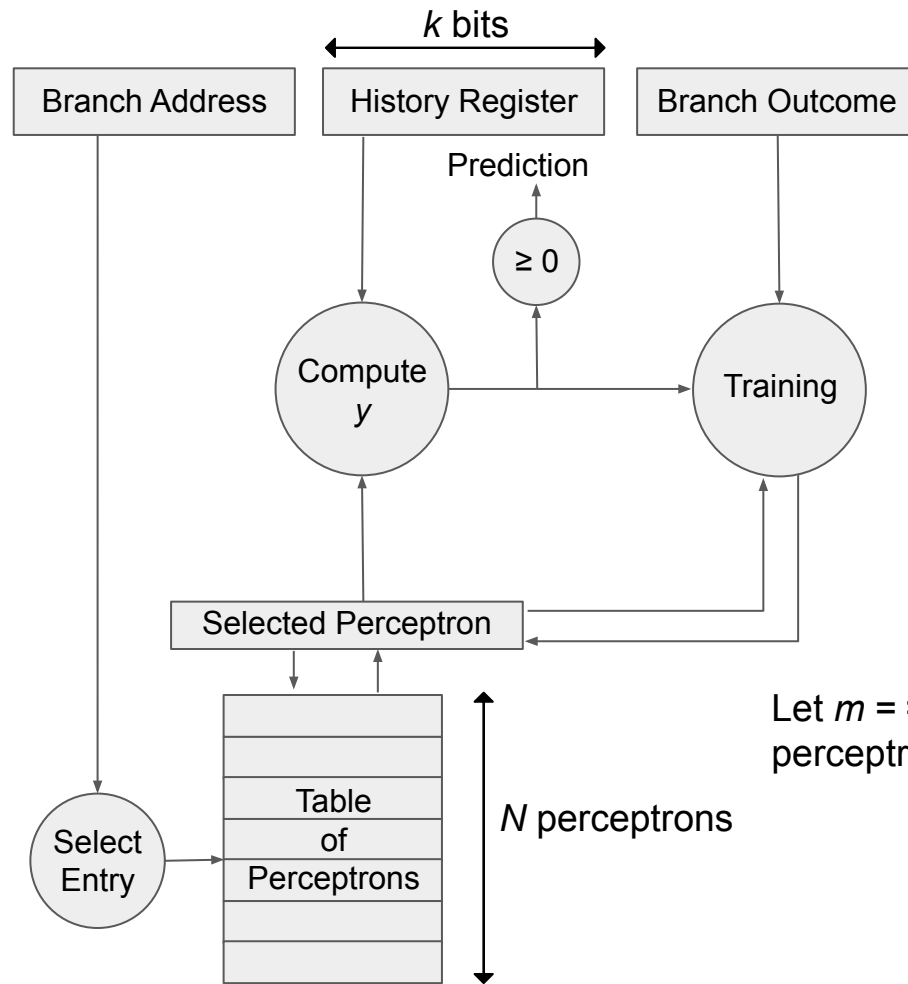




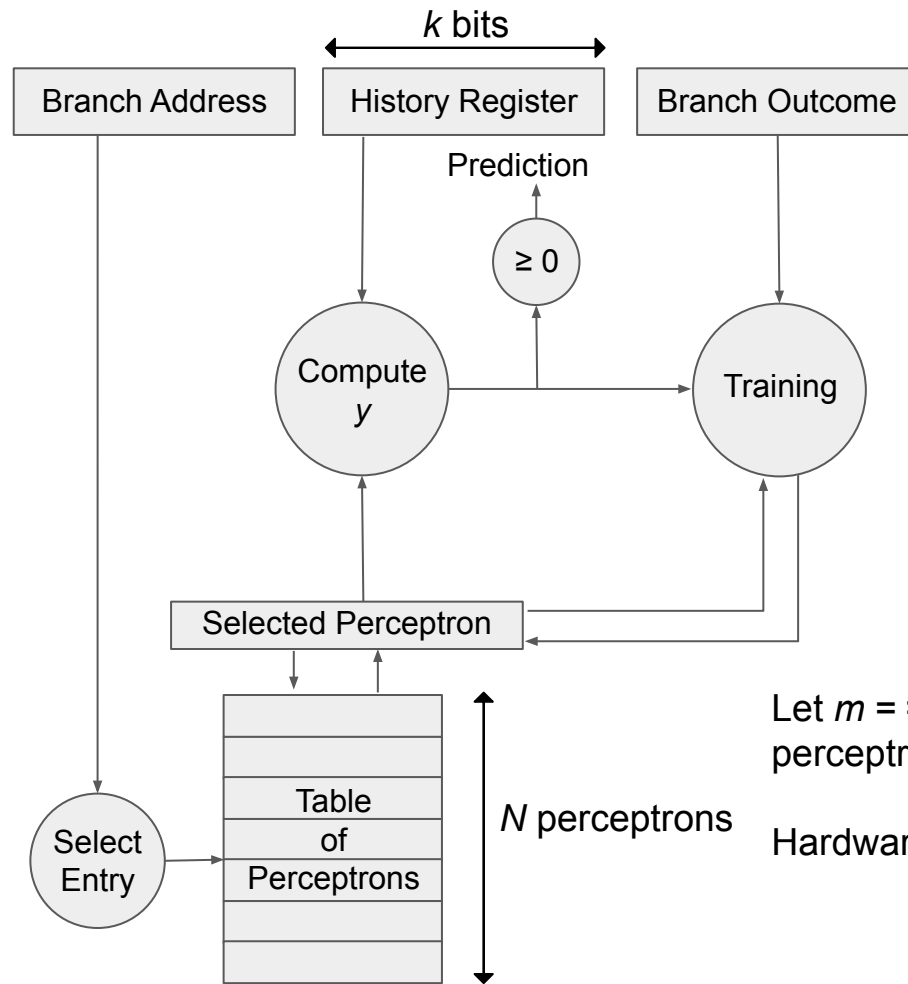








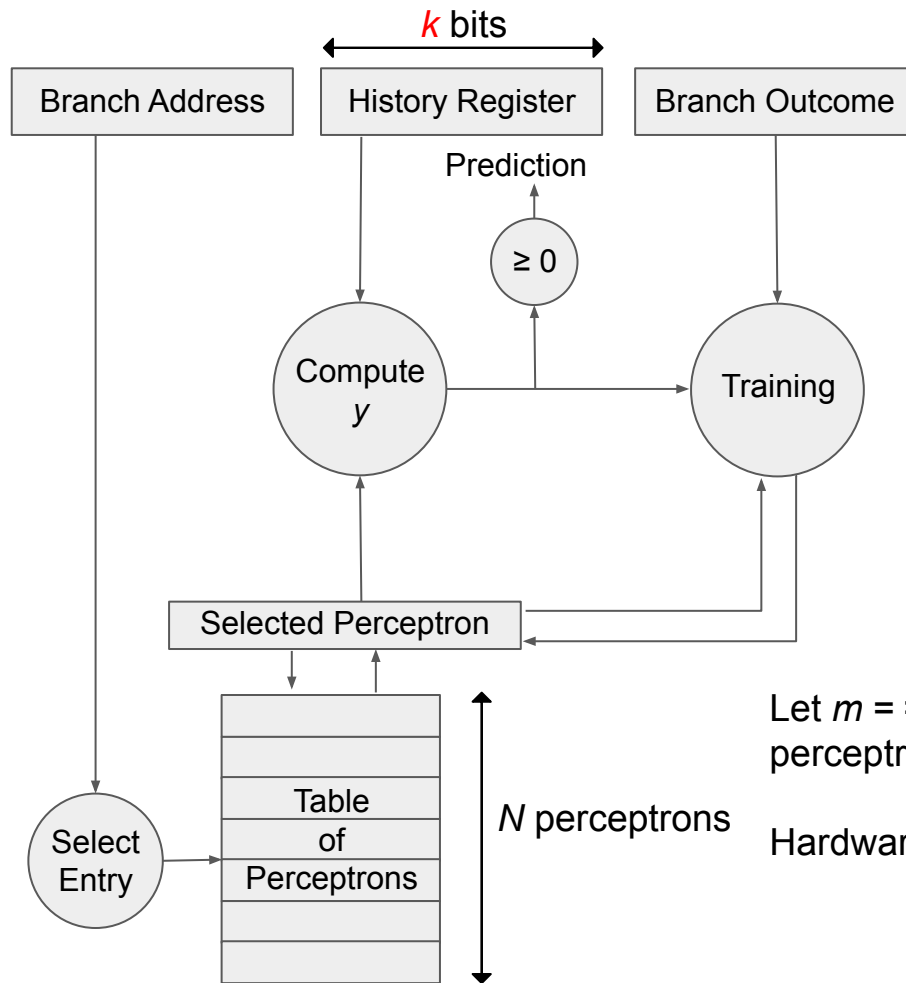
Let m = # of bits for each weight in a perceptron



Let $m = \#$ of bits for each weight in a perceptron

Hardware resources = $N * (k + 1) * m$

Hardware resources scale
linearly with the history
length!

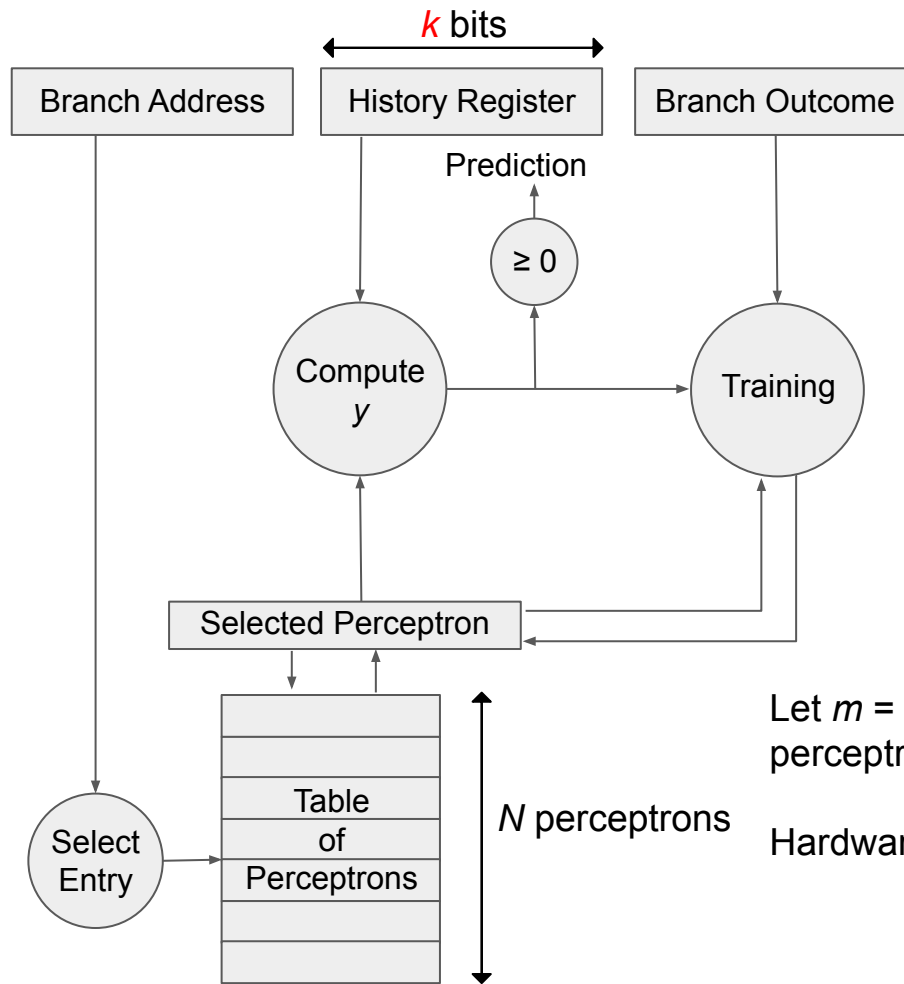


Let m = # of bits for each weight in a perceptron

Hardware resources = $N * (k + 1) * m$

Hardware resources scale **linearly** with the history length!

Can consider much longer history lengths than a two-level predictor



Let m = # of bits for each weight in a perceptron

Hardware resources = $N * (k + 1) * m$

By leveraging longer history lengths, neural branch predictors achieve higher prediction accuracy

Problem

$$y = \sum_{i=0}^n x_i w_i$$

Computing a dot product each time you make a prediction is too slow!

Solutions

Solutions

- Avoid doing multiplication since each x_i is 1 or -1
 - Only need to add/subtract weights to compute y

Solutions

- Avoid doing multiplication since each x_i is 1 or -1
 - Only need to add/subtract weights to compute y
- Only need the most significant bit of y to make a prediction
 - Make prediction before all of y is computed

Solutions

- Avoid doing multiplication since each x_i is 1 or -1
 - Only need to add/subtract weights to compute y
- Only need the most significant bit of y to make a prediction
 - Make prediction before all of y is computed
- Override predictors
 - Have 2 predictors, a faster less accurate predictor and a slower more accurate predictor
 - If slower predictor disagrees with faster predictor, override prediction

Solutions

Fast Path-Based Neural Branch Prediction

Daniel A. Jiménez

Department of Computer Science

Rutgers University, Piscataway, NJ 08854

Solutions

Fast Path-Based Neural Branch Prediction

Daniel A. Jiménez

Department of Computer Science

Rutgers University, Piscataway, NJ 08854

Staggers computation in time

Solutions

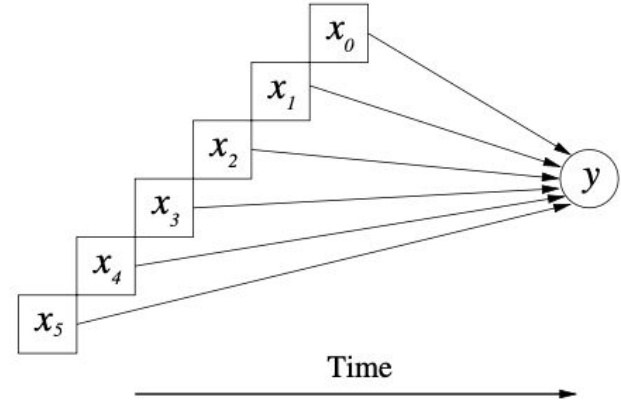
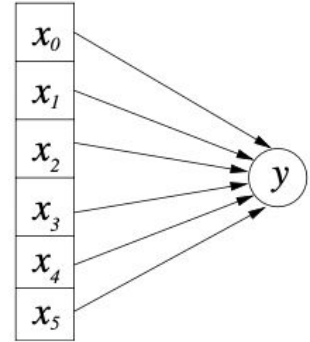
Fast Path-Based Neural Branch Prediction

Daniel A. Jiménez

Department of Computer Science
Rutgers University, Piscataway, NJ 08854

Staggeres computation in time

Predicts a branch using neurons selected dynamically along the path to the branch, rather than selecting the neurons all at once based solely on the branch address



Solutions

Fast Path-Based Neural Branch Prediction

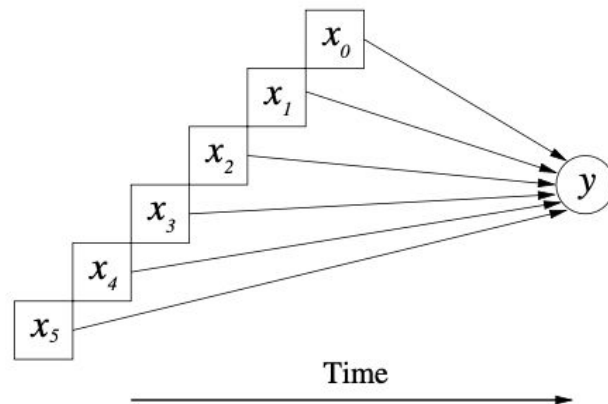
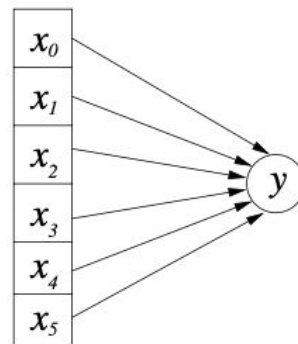
Daniel A. Jiménez

Department of Computer Science
Rutgers University, Piscataway, NJ 08854

Staggeres computation in time

Predicts a branch using neurons selected dynamically along the path to the branch, rather than selecting the neurons all at once based solely on the branch address

Faster and more accurate



Linear Separability

Linear Separability

Perceptrons can only learn functions that are linearly separable

Linear Separability

Perceptrons can only learn functions that are linearly separable

- A boolean function is linearly separable if and only if there exists a hyperplane that separates all of the true instances from all of the false instances

Linear Separability

Perceptrons can only learn functions that are linearly separable

- A boolean function is linearly separable if and only if there exists a hyperplane that separates all of the true instances from all of the false instances

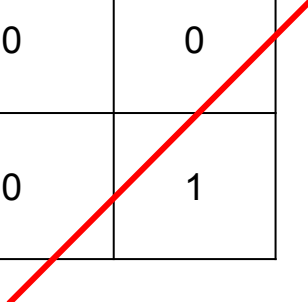
AND	0	1
0	0	0
1	0	1

Linear Separability

Perceptrons can only learn functions that are linearly separable

- A boolean function is linearly separable if and only if there exists a hyperplane that separates all of the true instances from all of the false instances

AND	0	1
0	0	0
1	0	1

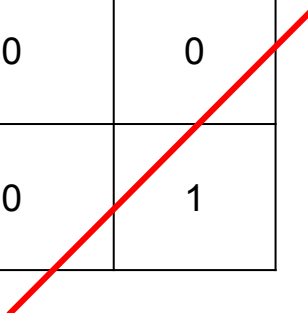


Linear Separability

Perceptrons can only learn functions that are linearly separable

- A boolean function is linearly separable if and only if there exists a hyperplane that separates all of the true instances from all of the false instances

AND	0	1
0	0	0
1	0	1



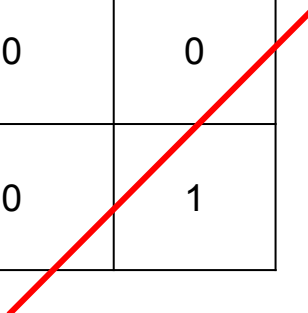
XOR	0	1
0	0	1
1	1	0

Linear Separability

Perceptrons can only learn functions that are linearly separable

- A boolean function is linearly separable if and only if there exists a hyperplane that separates all of the true instances from all of the false instances
- Most branch correlations are linearly separable

AND	0	1
0	0	0
1	0	1



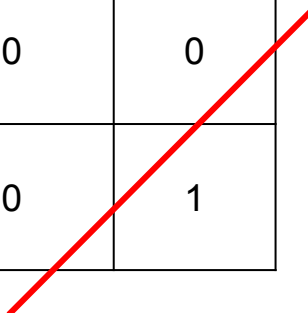
XOR	0	1
0	0	1
1	1	0

Linear Separability

Perceptrons can only learn functions that are linearly separable

- A boolean function is linearly separable if and only if there exists a hyperplane that separates all of the true instances from all of the false instances
- Most branch correlations are linearly separable
- Neural predictors work well in hybrid predictors

AND	0	1
0	0	0
1	0	1



XOR	0	1
0	0	1
1	1	0

Have architectures implemented neural
predictors?



Yes



Dynamic Branch Prediction with Perceptrons

Daniel A. Jiménez Calvin Lin

Department of Computer Sciences

The University of Texas at Austin

Austin, TX 78712

`{ djimenez, lin}@cs.utexas.edu`



Dynamic Branch Prediction with Perceptrons

Daniel A. Jiménez Calvin Lin

Department of Computer Sciences

The University of Texas at Austin

Austin, TX 78712

`{ djimenez, lin}@cs.utexas.edu`



2019 HPCA Test of Time Award

The HPCA Test of Time (ToT) award recognizes the most influential papers published in prior sessions of the International Symposium of High Performance Computer Architecture (HPCA) each of whom have had significant impact in the field.

Dynamic Branch Prediction with Perceptrons

Daniel A. Jiménez Calvin Lin

Department of Computer Sciences

The University of Texas at Austin

Austin, TX 78712

`{ djimenez, lin}@cs.utexas.edu`



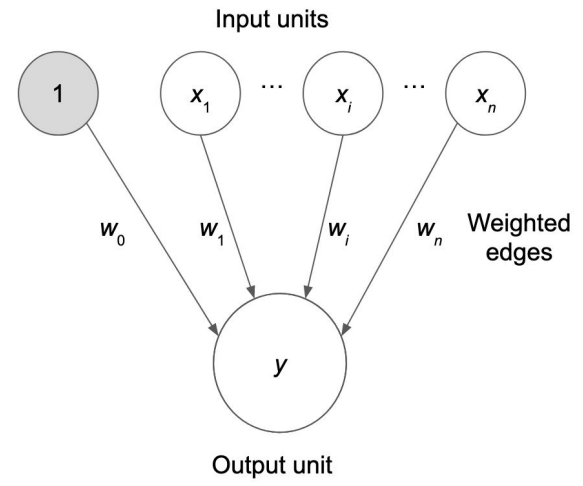
2019 HPCA Test of Time Award

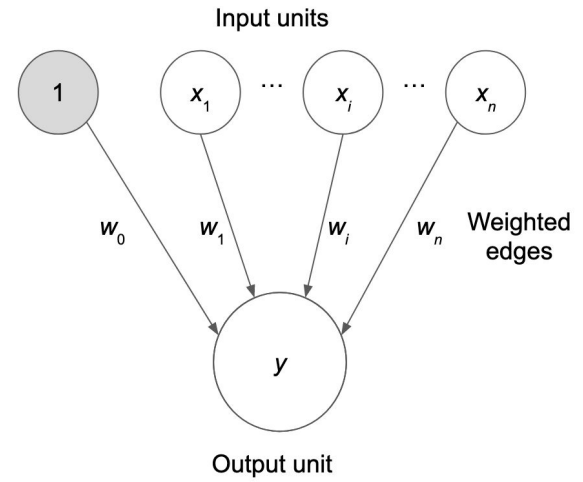
The HPCA Test of Time (ToT) award recognizes the most influential papers published in prior sessions of the International Symposium of High Performance Computer Architecture (HPCA) each of whom have had significant impact in the field.

2021 B. Ramakrishna Rau Award

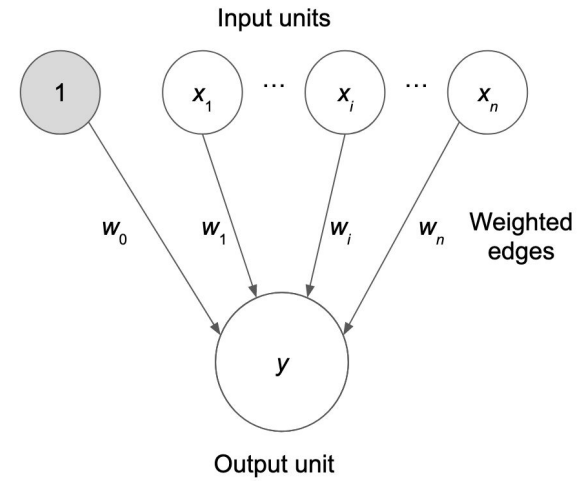
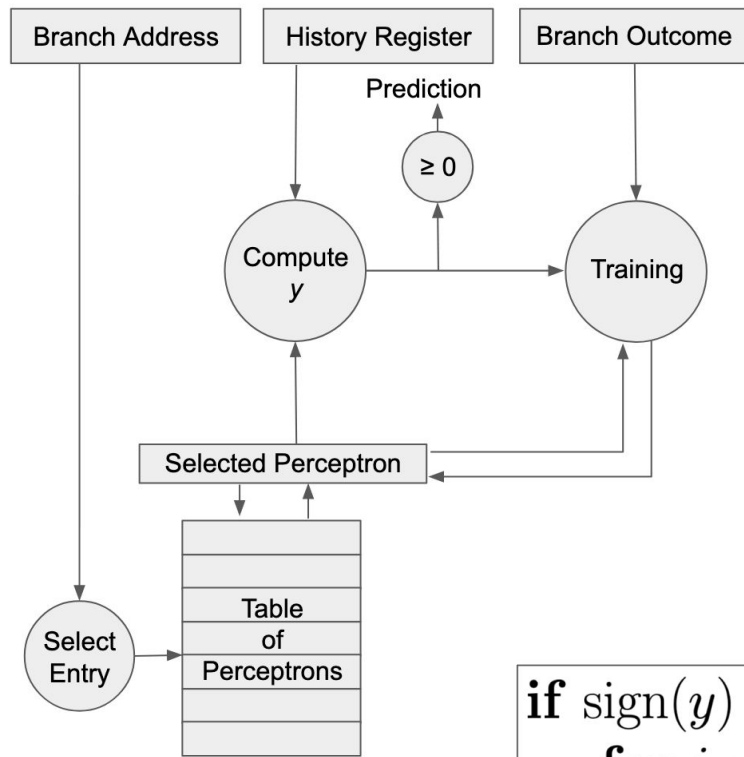
The B. Ramakrishna Rau award will be presented ***“in recognition of substantial contributions in the field of computer microarchitecture and compiler code generation.”***



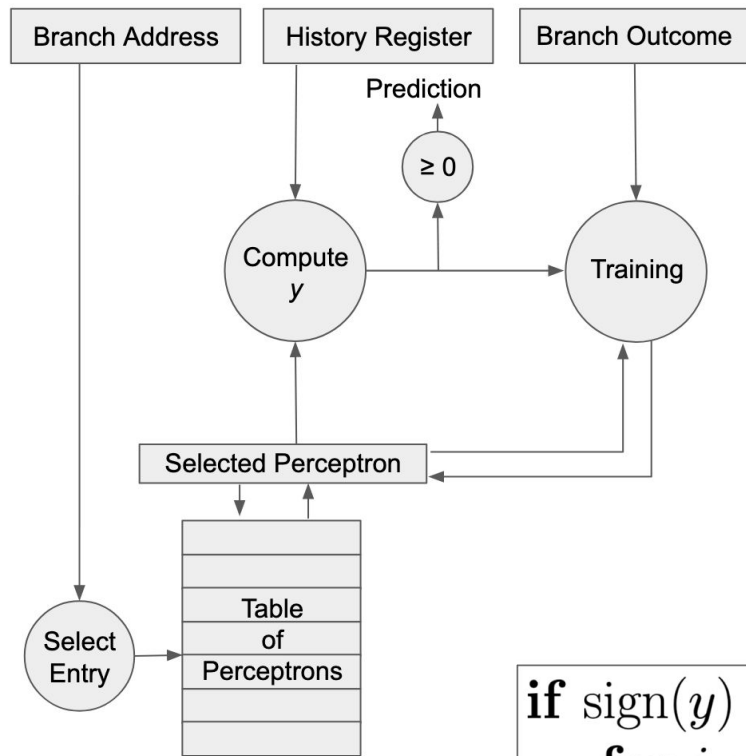




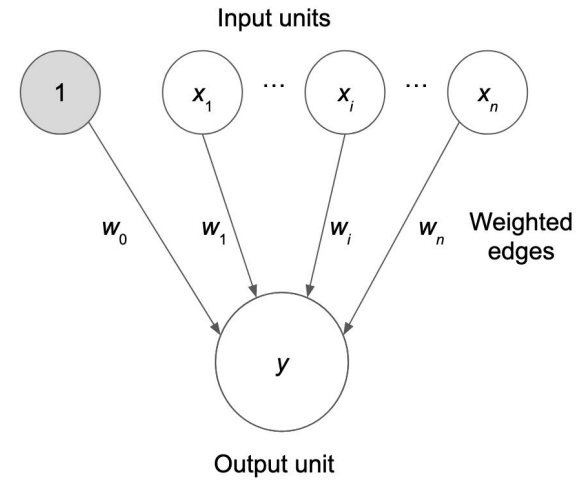
```
if  $\text{sign}(y) \neq t$  or  $|y| \leq \theta$  then  
  for  $i := 0$  to  $n$  do  
     $w_i := w_i + tx_i$   
  end for  
end if
```



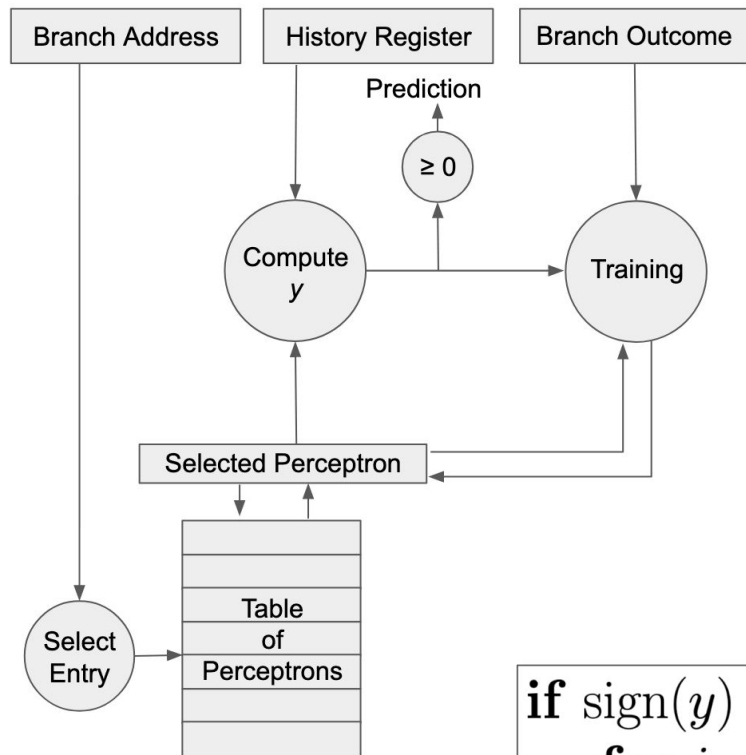
if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**
 for $i := 0$ to n **do**
 $w_i := w_i + tx_i$
 end for
end if



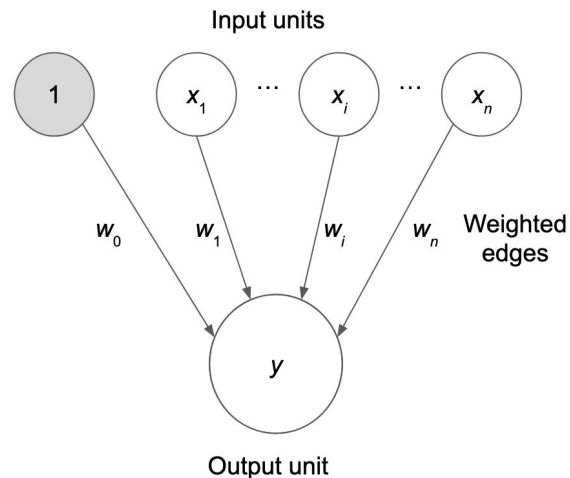
exponential vs. **linear**
scaling of hardware
resources



if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**
 for $i := 0$ to n **do**
 $w_i := w_i + tx_i$
 end for
end if

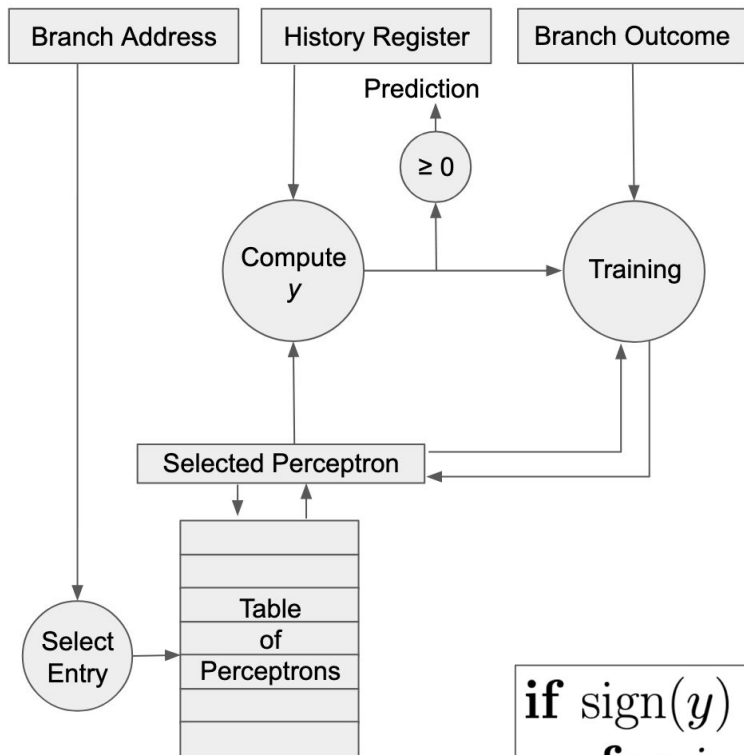


exponential vs. **linear**
scaling of hardware
resources



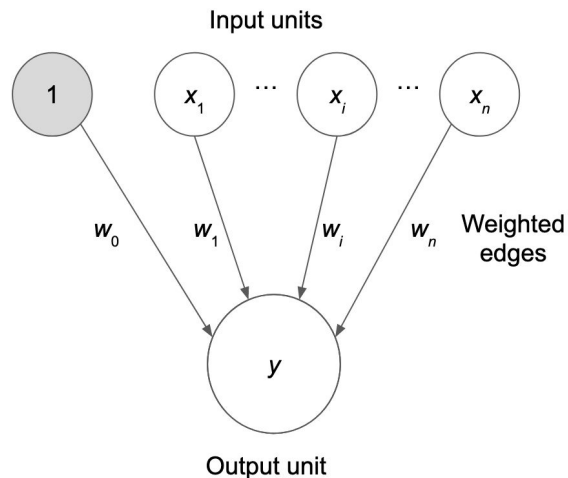
Neural branch predictors
leverage longer branch histories

if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**
 for $i := 0$ to n **do**
 $w_i := w_i + tx_i$
 end for
end if

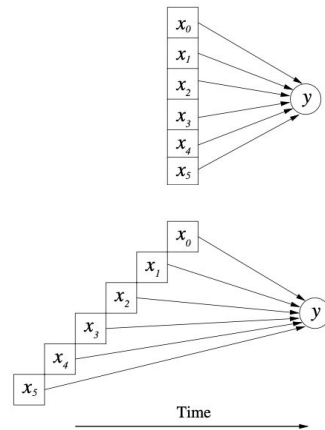


exponential vs. **linear**
scaling of hardware
resources

if $\text{sign}(y) \neq t$ or $|y| \leq \theta$ **then**
 for $i := 0$ to n **do**
 $w_i := w_i + tx_i$
 end for
end if



Neural branch predictors
leverage longer branch histories



Thank you :)

Have a great
reading week!

References

D. A. Jimenez and C. Lin, "Dynamic branch prediction with perceptrons," *Proceedings HPCA Seventh International Symposium on High-Performance Computer Architecture*, Monterrey, Mexico, 2001, pp. 197-206, doi: 10.1109/HPCA.2001.903263.

Daniel A. Jiménez and Calvin Lin. 2002. Neural methods for dynamic branch prediction. *ACM Trans. Comput. Syst.* 20, 4 (November 2002), 369–397. <https://doi.org/10.1145/571637.571639>

D. A. Jimenez, "Fast path-based neural branch prediction," *Proceedings. 36th Annual IEEE/ACM International Symposium on Microarchitecture, 2003. MICRO-36.*, San Diego, CA, USA, 2003, pp. 243-252, doi: 10.1109/MICRO.2003.1253199.